

```
1 from google.colab import drive
2 drive.mount('/content/drive')

→ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_

1 import numpy as np
2 import os
3 import tensorflow as tf
4 from tensorflow.keras.preprocessing.image import ImageDataGenerator
5 from tensorflow.keras.applications import DenseNet121
6 from tensorflow.keras.layers import Dropout, Dense, GlobalAveragePooling2D, Conv2D, MaxPooling2D, Flatten, BatchNor
7 from tensorflow.keras.models import Model, Sequential
8 from sklearn.metrics import classification_report, accuracy_score
9 from sklearn.utils import compute_class_weight
10 import matplotlib.pyplot as plt
11 from sklearn.metrics import confusion_matrix
12 import seaborn as sns
13 from imblearn.over_sampling import SMOTE
14
15
16
17
18 base_dir = "/content/drive/MyDrive/TB_Chest_Radiography_Database"
19
20
21 train_datagen = ImageDataGenerator(
22     rescale=1.0 / 255.0,
23     rotation_range=20,
24     zoom_range=0.2,
25     width_shift_range=0.2,
26     height_shift_range=0.2,
27     shear_range=0.2,
28     brightness_range=[0.8, 1.2],
29     horizontal_flip=True,
30     validation_split=0.2
```

```
31 )
32
33 val_datagen = ImageDataGenerator(rescale=1.0 / 255.0, validation_split=0.2)
34
35
36 train_data = train_datagen.flow_from_directory(
37     base_dir,
38     target_size=(224, 224),
39     batch_size=32,
40     class_mode="binary",
41     subset="training"
42 )
43
44 val_data = val_datagen.flow_from_directory(
45     base_dir,
46     target_size=(224, 224),
47     batch_size=32,
48     class_mode="binary",
49     shuffle=False,
50     subset="validation"
51 )
52
53
54 class_weights = compute_class_weight('balanced', classes=np.unique(train_data.classes), y=train_data.classes)
55 class_weights = {i: class_weights[i] for i in range(len(class_weights))}

56
57
58 def visualize_samples(data_generator, title):
59     images, labels = next(data_generator)
60     plt.figure(figsize=(10, 10))
61     for i in range(9):
62         plt.subplot(3, 3, i + 1)
63         plt.imshow(images[i])
64         plt.title(f"Label: {int(labels[i])}")
65         plt.axis("off")
66     plt.suptitle(title, fontsize=16)
67     plt.show()
```

```
//      plt.show()
68
69 print("Visualizing some training samples after augmentation...")
70 visualize_samples(train_data, "Training Samples with Augmentation")
71
72 print("Visualizing some validation samples...")
73 visualize_samples(val_data, "Validation Samples")
74
75
76 train_classes = train_data.classes
77 val_classes = val_data.classes
78
79 def plot_class_distribution(classes, title):
80     unique, counts = np.unique(classes, return_counts=True)
81     plt.figure(figsize=(8, 5))
82     sns.barplot(x=unique, y=counts, palette="viridis")
83     plt.title(title, fontsize=14)
84     plt.xlabel("Class", fontsize=12)
85     plt.ylabel("Number of Samples", fontsize=12)
86     plt.xticks(ticks=unique, labels=["Normal", "Tuberculosis"])
87     plt.show()
88     return unique, counts
89
90 print("Class distribution BEFORE handling imbalance:")
91 train_unique, train_counts = plot_class_distribution(train_classes, "Training Class Distribution (Before Balancing)")
92 val_unique, val_counts = plot_class_distribution(val_classes, "Validation Class Distribution")
93
94
95 X_train = []
96 y_train = []
97 for _ in range(len(train_data)):
98     images, labels = next(train_data)
99     X_train.append(images)
100    y_train.append(labels)
101
102 X_train = np.concatenate(X_train, axis=0).reshape(-1, 224 * 224 * 3)
103 y_train = np.concatenate(y_train, axis=0)
104
```

```
104  
105  
106  
107 smote = SMOTE(random_state=42)  
108 X_train_res, y_train_res = smote.fit_resample(X_train, y_train)  
109  
110  
111 X_train_res_images = X_train_res.reshape(-1, 224, 224, 3)  
112  
113  
114 print("\nClass distribution AFTER handling imbalance with SMOTE:")  
115 train_res_unique, train_res_counts = plot_class_distribution(y_train_res, "Training Class Distribution (After Balar  
116  
117  
118 def visualize_smote_samples(X_res, y_res, title):  
119     plt.figure(figsize=(10, 10))  
120     for i in range(9):  
121         plt.subplot(3, 3, i + 1)  
122         plt.imshow(X_res[i])  
123         plt.title(f"Label: {int(y_res[i])}")  
124         plt.axis("off")  
125     plt.suptitle(title, fontsize=16)  
126     plt.show()  
127  
128 print("Visualizing some synthetic samples after SMOTE...")  
129 visualize_smote_samples(X_train_res_images, y_train_res, "Synthetic Samples After SMOTE")  
130
```

Found 3360 images belonging to 2 classes.

Found 840 images belonging to 2 classes.

Visualizing some training samples after augmentation...

Training Samples with Augmentation

Label: 1



Label: 0



Label: 0





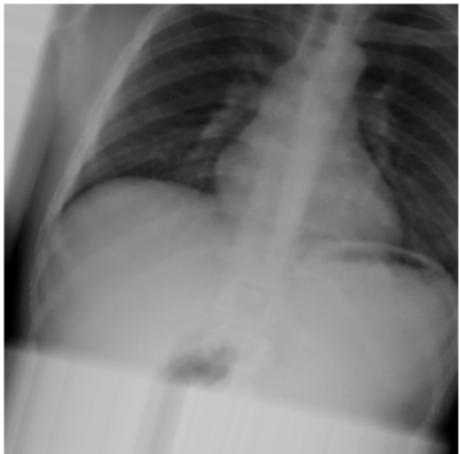
Label: 0



Label: 1



Label: 1



Label: 0



Label: 1



Label: 1

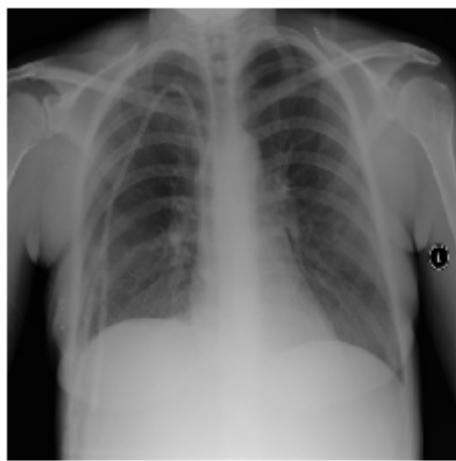




Visualizing some validation samples...

Validation Samples

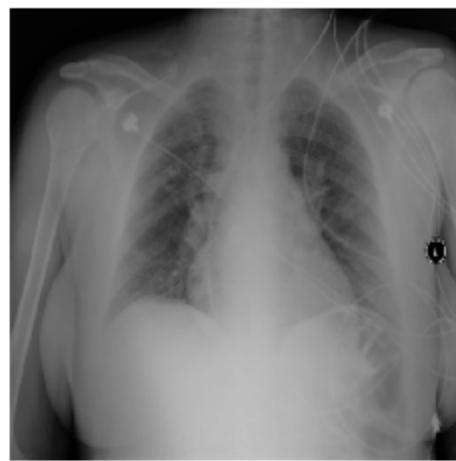
Label: 0



Label: 0



Label: 0



Label: 0

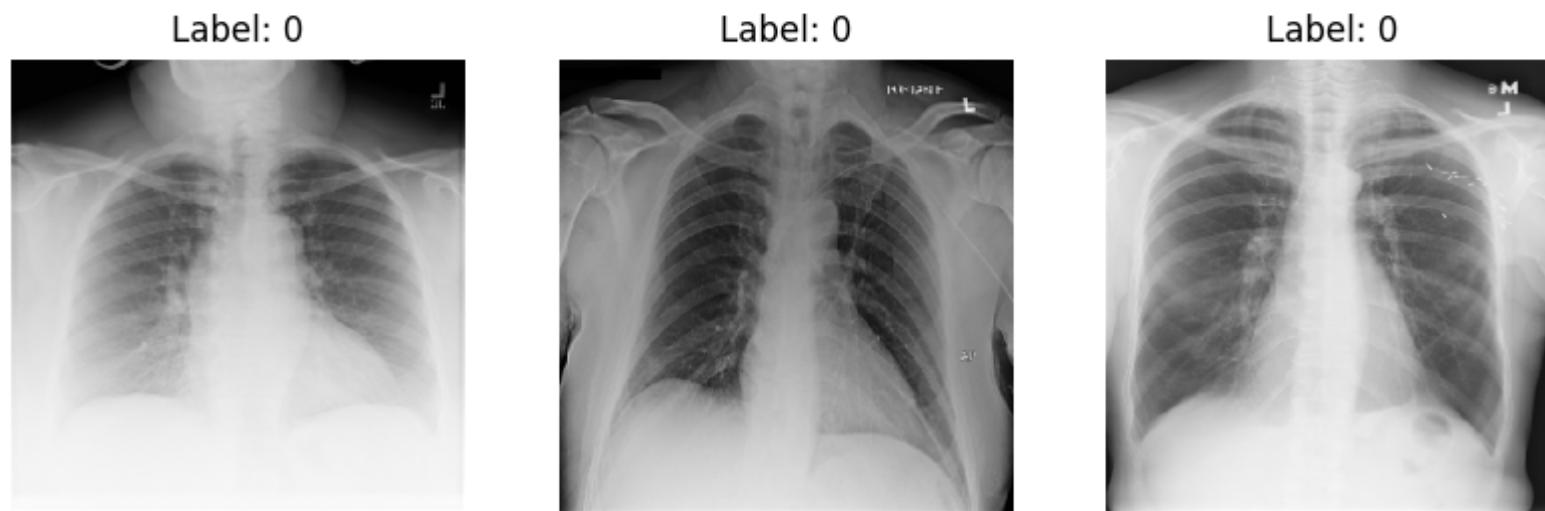


Label: 0



Label: 0





Class distribution BEFORE handling imbalance:

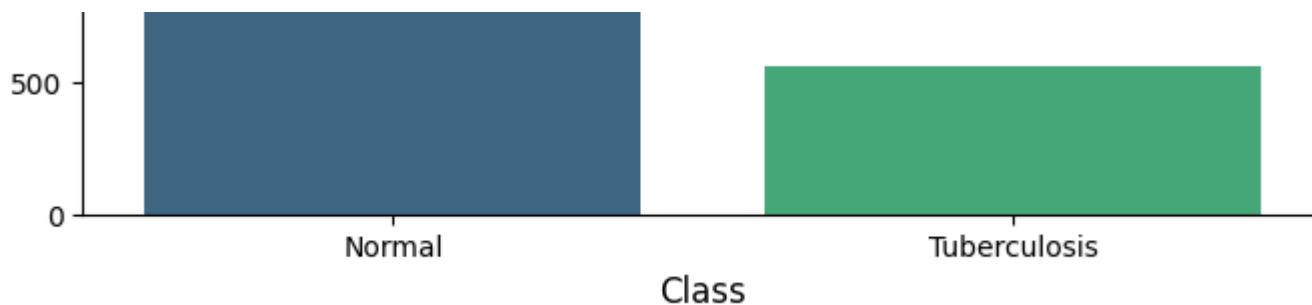
```
<ipython-input-3-7349cddcffdc>:82: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to

```
sns.barplot(x=unique, y=counts, palette="viridis")
```

Training Class Distribution (Before Balancing)



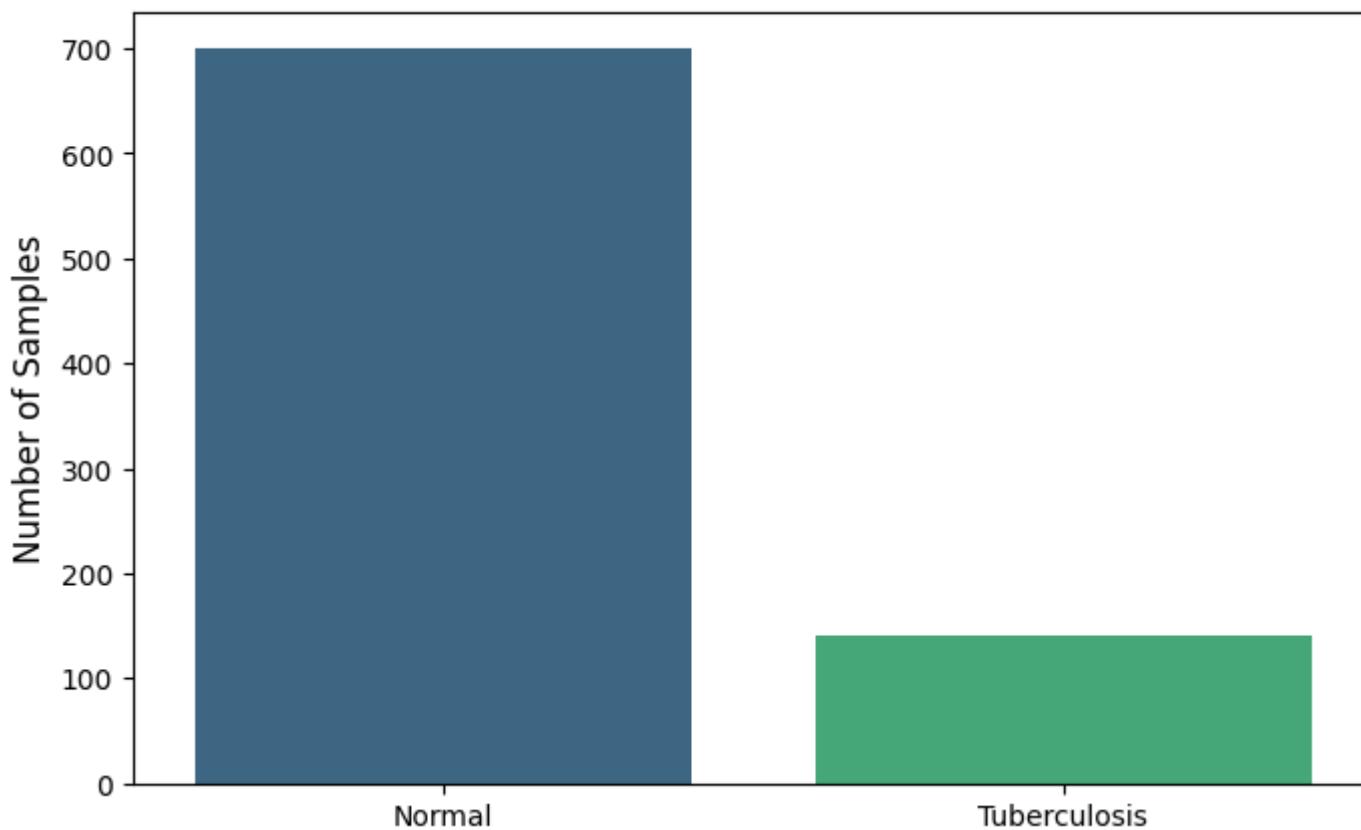


```
<ipython-input-3-7349cddcffdc>:82: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to
```

```
sns.barplot(x=unique, y=counts, palette="viridis")
```

Validation Class Distribution



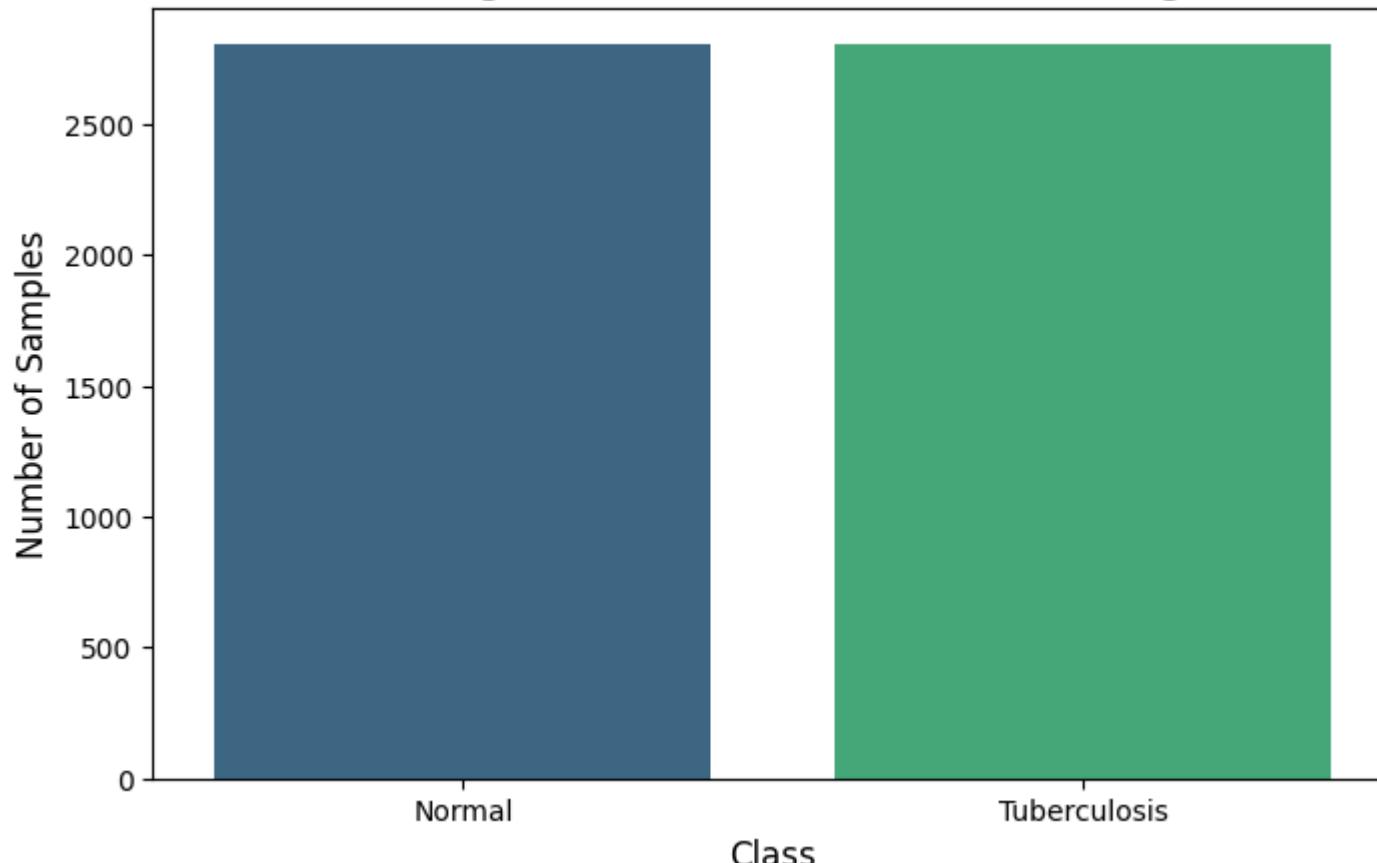
CLASS

Class distribution AFTER handling imbalance with SMOTE:
<ipython-input-3-7349cddcffdc>:82: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to

```
sns.barplot(x=unique, y=counts, palette="viridis")
```

Training Class Distribution (After Balancing)



Visualizing some synthetic samples after SMOTE...

Synthetic Samples After SMOTE

Label: 0



Label: 0



Label: 0



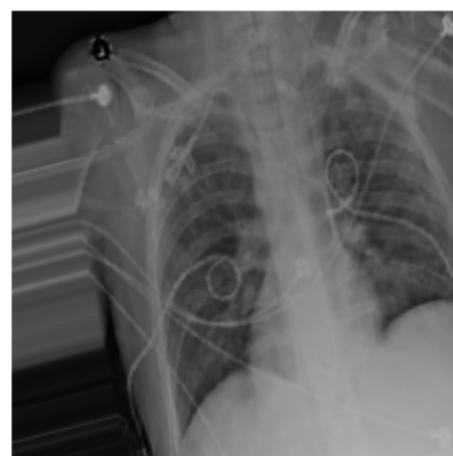
Label: 0



Label: 0



Label: 0



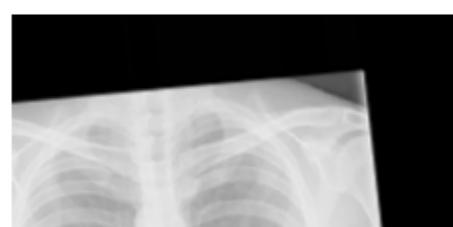
Label: 0



Label: 0



Label: 1





```
1
2 cnn_model = Sequential([
3     Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224, 3)),
4     MaxPooling2D((2, 2)),
5     BatchNormalization(),
6
7     Conv2D(64, (3, 3), activation='relu'),
8     MaxPooling2D((2, 2)),
9     BatchNormalization(),
10
11    Flatten(),
12    Dense(64, kernel_initializer='uniform', activation='relu'),
13    Dropout(0.4),
14
15    Dense(1, activation='sigmoid')
16 ])
17
18 cnn_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
19 cnn_history = cnn_model.fit(X_train_res_images, y_train_res, epochs=20, class_weight=class_weights)
20
21 cnn_predictions = cnn_model.predict(val_data, steps=int(np.ceil(val_data.samples / val_data.batch_size))), verbose=2
22 cnn_predictions = (cnn_predictions > 0.6).astype(int)
23
24 print("CNN Model Accuracy:", accuracy_score(val_data.classes, cnn_predictions))
25 print("CNN Classification Report:\n", classification_report(val_data.classes, cnn_predictions))

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/20
/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:122: UserWarning: You
  self._warn_if_super_not_called()
105/105 418s 4s/step - accuracy: 0.7704 - loss: 20.6091 - val_accuracy: 0.1679 - val_loss: 4.
Epoch 2/20
105/105 117s 798ms/step - accuracy: 0.7886 - loss: 5.4675 - val_accuracy: 0.1810 - val_loss:
Epoch 3/20
105/105 143s 834ms/step - accuracy: 0.8652 - loss: 1.6285 - val_accuracy: 0.7798 - val_loss:
Epoch 4/20
105/105 99s 837ms/step - accuracy: 0.8651 - loss: 2.9271 - val_accuracy: 0.8833 - val_loss: 6
Epoch 5/20
105/105 134s 825ms/step - accuracy: 0.8959 - loss: 2.0373 - val_accuracy: 0.9357 - val_loss:
Epoch 6/20
105/105 144s 832ms/step - accuracy: 0.9046 - loss: 1.6156 - val_accuracy: 0.1667 - val_loss:
Epoch 7/20
105/105 142s 830ms/step - accuracy: 0.8759 - loss: 3.1349 - val_accuracy: 0.5988 - val_loss:
Epoch 8/20
105/105 95s 860ms/step - accuracy: 0.9064 - loss: 0.8253 - val_accuracy: 0.1690 - val_loss: 4
Epoch 9/20
105/105 134s 784ms/step - accuracy: 0.9140 - loss: 0.4592 - val_accuracy: 0.9357 - val_loss:
Epoch 10/20
105/105 89s 795ms/step - accuracy: 0.9293 - loss: 0.3737 - val_accuracy: 0.9345 - val_loss: 6
Epoch 11/20
105/105 89s 805ms/step - accuracy: 0.9297 - loss: 0.4088 - val_accuracy: 0.8738 - val_loss: 6
Epoch 12/20
105/105 91s 822ms/step - accuracy: 0.9157 - loss: 0.4597 - val_accuracy: 0.8905 - val_loss: 6
Epoch 13/20
105/105 87s 787ms/step - accuracy: 0.9268 - loss: 0.3636 - val_accuracy: 0.8690 - val_loss: 6
Epoch 14/20
105/105 146s 813ms/step - accuracy: 0.9309 - loss: 0.4449 - val_accuracy: 0.6988 - val_loss:
Epoch 15/20
105/105 145s 852ms/step - accuracy: 0.9153 - loss: 0.5416 - val_accuracy: 0.9357 - val_loss:
Epoch 16/20
105/105 141s 835ms/step - accuracy: 0.9203 - loss: 0.4118 - val_accuracy: 0.9333 - val_loss:
Epoch 17/20
105/105 142s 839ms/step - accuracy: 0.9165 - loss: 0.4041 - val_accuracy: 0.9464 - val_loss:
Epoch 18/20
105/105 92s 829ms/step - accuracy: 0.9307 - loss: 0.3270 - val_accuracy: 0.8714 - val_loss: 6
Epoch 19/20
```

105/105 ————— **144s** 836ms/step - accuracy: 0.9066 - loss: 0.6851 - val_accuracy: 0.8869 - val_loss:
Epoch 20/20
105/105 ————— **139s** 825ms/step - accuracy: 0.9256 - loss: 0.3783 - val_accuracy: 0.9310 - val_loss:
27/27 - 10s - 361ms/step
CNN Model Accuracy: 0.9178571428571428
CNN Classification Report:

	precision	recall	f1-score	support
0	0.91	1.00	0.95	700
1	0.99	0.51	0.68	140
accuracy			0.92	840
macro avg	0.95	0.76	0.81	840
weighted avg	0.92	0.92	0.91	840

```
1
2 densenet_model = DenseNet121(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
3 densenet_model.trainable = False
4
5 x = densenet_model.output
6 x = GlobalAveragePooling2D()(x)
7 dropout_layer = Dropout(0.5)(x)
8 output = Dense(1, activation='sigmoid')(dropout_layer)
9
10 final_densenet_model = Model(inputs=densenet_model.input, outputs=output)
11
12 final_densenet_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
13 densenet_history = final_densenet_model.fit(
14     X_train_res_images, y_train_res, epochs=20, class_weight=class_weights
15 )
16
17 densenet_predictions = final_densenet_model.predict(
18     val_data, steps=int(np.ceil(val_data.samples / val_data.batch_size)), verbose=2
19 )
20 densenet_predictions = (densenet_predictions > 0.6).astype(int)
21
22 print("DenseNet Model Accuracy: " accuracy_score(val_data.classes, densenet_predictions))
```

```
22 print("DenseNet Model Accuracy: ", accuracy_score(val_data.classes, densenet_predictions))
23 print("DenseNet Classification Report:\n", classification_report(val_data.classes, densenet_predictions))

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/densenet/densenet121\_weights\_tf-29084464/29084464
29084464/29084464 0s 0us/step
Epoch 1/20
105/105 144s 1s/step - accuracy: 0.5987 - loss: 0.7670 - val_accuracy: 0.8750 - val_loss: 0.3
Epoch 2/20
105/105 108s 849ms/step - accuracy: 0.8151 - loss: 0.4141 - val_accuracy: 0.9024 - val_loss:
Epoch 3/20
105/105 96s 859ms/step - accuracy: 0.8574 - loss: 0.3510 - val_accuracy: 0.9167 - val_loss: 0
Epoch 4/20
105/105 92s 835ms/step - accuracy: 0.8662 - loss: 0.3233 - val_accuracy: 0.9143 - val_loss: 0
Epoch 5/20
105/105 90s 804ms/step - accuracy: 0.8836 - loss: 0.2975 - val_accuracy: 0.9250 - val_loss: 0
Epoch 6/20
105/105 153s 927ms/step - accuracy: 0.9039 - loss: 0.2478 - val_accuracy: 0.9333 - val_loss:
Epoch 7/20
105/105 91s 811ms/step - accuracy: 0.9048 - loss: 0.2372 - val_accuracy: 0.9250 - val_loss: 0
Epoch 8/20
105/105 145s 851ms/step - accuracy: 0.9092 - loss: 0.2515 - val_accuracy: 0.9214 - val_loss:
Epoch 9/20
105/105 93s 835ms/step - accuracy: 0.9144 - loss: 0.2271 - val_accuracy: 0.9274 - val_loss: 0
Epoch 10/20
105/105 144s 853ms/step - accuracy: 0.9158 - loss: 0.2150 - val_accuracy: 0.9286 - val_loss:
Epoch 11/20
105/105 91s 824ms/step - accuracy: 0.9262 - loss: 0.2152 - val_accuracy: 0.9286 - val_loss: 0
Epoch 12/20
105/105 92s 821ms/step - accuracy: 0.9140 - loss: 0.2348 - val_accuracy: 0.9286 - val_loss: 0
Epoch 13/20
105/105 92s 832ms/step - accuracy: 0.9200 - loss: 0.2230 - val_accuracy: 0.9345 - val_loss: 0
Epoch 14/20
105/105 94s 843ms/step - accuracy: 0.9109 - loss: 0.2295 - val_accuracy: 0.9321 - val_loss: 0
Epoch 15/20
105/105 139s 826ms/step - accuracy: 0.9086 - loss: 0.2302 - val_accuracy: 0.9512 - val_loss:
Epoch 16/20
105/105 90s 801ms/step - accuracy: 0.9009 - loss: 0.2504 - val_accuracy: 0.9440 - val_loss: 0
Epoch 17/20
105/105 144s 838ms/step - accuracy: 0.9269 - loss: 0.1995 - val_accuracy: 0.9548 - val_loss:
Epoch 18/20
105/105 91s 814ms/step - accuracy: 0.9152 - loss: 0.2127 - val_accuracy: 0.9607 - val_loss: 0
```

```
105/105 ━━━━━━━━━━ 115 0.14ms/step - accuracy: 0.9102 - loss: 0.2121 - val_accuracy: 0.9007 - val_loss: 0
Epoch 19/20
105/105 ━━━━━━━━━━ 141s 822ms/step - accuracy: 0.9211 - loss: 0.2065 - val_accuracy: 0.9607 - val_loss:
Epoch 20/20
105/105 ━━━━━━━━━━ 143s 823ms/step - accuracy: 0.9106 - loss: 0.2247 - val_accuracy: 0.9595 - val_loss:
27/27 - 27s - 1s/step
DenseNet Model Accuracy: 0.9464285714285714
```

DenseNet Classification Report:

	precision	recall	f1-score	support
0	0.95	0.99	0.97	700
1	0.94	0.72	0.82	140
accuracy			0.95	840
macro avg	0.95	0.86	0.89	840
weighted avg	0.95	0.95	0.94	840

```
1
2 from tensorflow.keras.applications import ResNet50
3 from tensorflow.keras.layers import GlobalAveragePooling2D, Dropout, Dense
4 from tensorflow.keras.models import Model
5
6 resnet_model = ResNet50(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
7 resnet_model.trainable = False
8
9 x = resnet_model.output
10 x = GlobalAveragePooling2D()(x)
11 dropout_layer = Dropout(0.5)(x)
12 output = Dense(1, activation='sigmoid')(dropout_layer)
13
14 final_resnet_model = Model(inputs=resnet_model.input, outputs=output)
15
16 final_resnet_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
17 resnet_history = final_resnet_model.fit(
18     X_train_res_images, y_train_res, epochs=20, class_weight=class_weights
19 )
20
```

```
21 resnet_predictions = final_resnet_model.predict(  
22     val_data, steps=int(np.ceil(val_data.samples / val_data.batch_size)), verbose=2  
23 )  
24 resnet_predictions = (resnet_predictions > 0.6).astype(int)  
25  
26 print("ResNet50 Model Accuracy:", accuracy_score(val_data.classes, resnet_predictions))  
27 print("ResNet50 Classification Report:\n", classification_report(val_data.classes, resnet_predictions))  
28
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_94765736/94765736 0s 0us/step

Epoch 1/20
105/105 117s 921ms/step - accuracy: 0.5390 - loss: 0.7437 - val_accuracy: 0.8333 - val_loss: 0
Epoch 2/20
105/105 96s 859ms/step - accuracy: 0.5603 - loss: 0.7079 - val_accuracy: 0.8333 - val_loss: 0
Epoch 3/20
105/105 94s 848ms/step - accuracy: 0.5844 - loss: 0.6934 - val_accuracy: 0.2321 - val_loss: 0
Epoch 4/20
105/105 97s 876ms/step - accuracy: 0.5879 - loss: 0.6764 - val_accuracy: 0.3226 - val_loss: 0
Epoch 5/20
105/105 95s 845ms/step - accuracy: 0.5594 - loss: 0.6824 - val_accuracy: 0.2643 - val_loss: 0
Epoch 6/20
105/105 141s 844ms/step - accuracy: 0.4973 - loss: 0.6951 - val_accuracy: 0.1667 - val_loss: 0
Epoch 7/20
105/105 142s 844ms/step - accuracy: 0.4628 - loss: 0.6908 - val_accuracy: 0.3119 - val_loss: 0
Epoch 8/20
105/105 144s 851ms/step - accuracy: 0.5746 - loss: 0.6816 - val_accuracy: 0.3976 - val_loss: 0
Epoch 9/20
105/105 137s 818ms/step - accuracy: 0.5534 - loss: 0.6898 - val_accuracy: 0.1667 - val_loss: 0
Epoch 10/20
105/105 92s 835ms/step - accuracy: 0.4976 - loss: 0.6899 - val_accuracy: 0.2821 - val_loss: 0
Epoch 11/20
105/105 92s 830ms/step - accuracy: 0.4596 - loss: 0.6874 - val_accuracy: 0.8262 - val_loss: 0
Epoch 12/20
105/105 142s 821ms/step - accuracy: 0.5827 - loss: 0.6935 - val_accuracy: 0.1714 - val_loss: 0
Epoch 13/20
105/105 90s 810ms/step - accuracy: 0.5520 - loss: 0.6827 - val_accuracy: 0.4774 - val_loss: 0
Epoch 14/20
105/105 93s 843ms/step - accuracy: 0.5558 - loss: 0.6694 - val_accuracy: 0.5452 - val_loss: 0
Epoch 15/20

```
Epoch 15/20  
105/105 92s 837ms/step - accuracy: 0.5240 - loss: 0.7033 - val_accuracy: 0.8083 - val_loss: 0  
Epoch 16/20  
105/105 96s 858ms/step - accuracy: 0.6696 - loss: 0.6685 - val_accuracy: 0.1667 - val_loss: 0  
Epoch 17/20  
105/105 94s 847ms/step - accuracy: 0.5150 - loss: 0.6738 - val_accuracy: 0.1667 - val_loss: 0  
Epoch 18/20  
105/105 94s 853ms/step - accuracy: 0.5376 - loss: 0.6777 - val_accuracy: 0.4988 - val_loss: 0  
Epoch 19/20  
105/105 95s 850ms/step - accuracy: 0.5959 - loss: 0.6760 - val_accuracy: 0.2679 - val_loss: 0  
Epoch 20/20  
105/105 141s 857ms/step - accuracy: 0.5792 - loss: 0.6871 - val_accuracy: 0.8190 - val_loss:  
27/27 - 18s - 657ms/step
```

BesNet50 Model Accuracy: 0.8333333333333334

ResNet50 Classification Report:

	precision	recall	f1-score	support
0	0.83	1.00	0.91	700
1	0.00	0.00	0.00	140
accuracy			0.83	840
macro avg	0.42	0.50	0.45	840
weighted avg	0.69	0.83	0.76	840

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
1 # EfficientNetB0 Model
2 from tensorflow.keras.applications import EfficientNetB0
3
4
5 # Load pre-trained EfficientNetB0
6 efficientnet_model = EfficientNetB0(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
7 efficientnet_model.trainable = False
8
9 x = efficientnet_model.output
10 x = GlobalAveragePooling2D()(x) # Pooling layer
```

```
11 dropout_layer = Dropout(0.5)(x) # Dropout to reduce overfitting
12 output = Dense(1, activation='sigmoid')(dropout_layer) # Output layer for binary classification
13
14 final_efficientnet_model = Model(inputs=efficientnet_model.input, outputs=output)
15
16 # Compile the model
17 final_efficientnet_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
18
19 # Train the model
20 efficientnet_history = final_efficientnet_model.fit(
21     X_train_res_images, y_train_res,
22     epochs=20,
23     class_weight=class_weights
24 )
25
26 # Predictions
27 efficientnet_predictions = final_efficientnet_model.predict(
28     val_data,
29     steps=int(np.ceil(val_data.samples / val_data.batch_size)),
30     verbose=2
31 )
32 efficientnet_predictions = (efficientnet_predictions > 0.6).astype(int)
33
34 # Evaluate the model
35 print("EfficientNetB0 Model Accuracy:", accuracy_score(val_data.classes, efficientnet_predictions))
36 print("EfficientNetB0 Classification Report:\n", classification_report(val_data.classes, efficientnet_predictions))
37
```

```
Epoch 1/20
/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:122: UserWarning: 
  self._warn_if_super_not_called()
105/105 666s 6s/step - accuracy: 0.5069 - loss: 0.6939 - val_accuracy: 0.1667 - val_loss: 0.7
Epoch 2/20
105/105 510s 5s/step - accuracy: 0.3460 - loss: 0.7048 - val_accuracy: 0.8333 - val_loss: 0.6
Epoch 3/20
105/105 546s 5s/step - accuracy: 0.5256 - loss: 0.6934 - val_accuracy: 0.8333 - val_loss: 0.6
Epoch 4/20
105/105 498s 5s/step - accuracy: 0.4665 - loss: 0.7021 - val_accuracy: 0.8333 - val_loss: 0.6
```

Epoch 5/20
105/105 487s 5s/step - accuracy: 0.5855 - loss: 0.6991 - val_accuracy: 0.8333 - val_loss: 0.6
Epoch 6/20
105/105 506s 5s/step - accuracy: 0.6224 - loss: 0.6843 - val_accuracy: 0.1667 - val_loss: 0.7
Epoch 7/20
105/105 492s 5s/step - accuracy: 0.3264 - loss: 0.6865 - val_accuracy: 0.1667 - val_loss: 0.7
Epoch 8/20
105/105 484s 5s/step - accuracy: 0.4171 - loss: 0.6827 - val_accuracy: 0.1667 - val_loss: 0.7
Epoch 9/20
105/105 487s 5s/step - accuracy: 0.3381 - loss: 0.6998 - val_accuracy: 0.1690 - val_loss: 0.6
Epoch 10/20
105/105 482s 5s/step - accuracy: 0.5498 - loss: 0.6909 - val_accuracy: 0.8333 - val_loss: 0.6
Epoch 11/20
105/105 561s 5s/step - accuracy: 0.7367 - loss: 0.6685 - val_accuracy: 0.1667 - val_loss: 0.7
Epoch 12/20
105/105 487s 5s/step - accuracy: 0.3213 - loss: 0.6925 - val_accuracy: 0.1667 - val_loss: 0.7
Epoch 13/20
105/105 482s 5s/step - accuracy: 0.3741 - loss: 0.6902 - val_accuracy: 0.8333 - val_loss: 0.6
Epoch 14/20
105/105 475s 4s/step - accuracy: 0.7207 - loss: 0.6829 - val_accuracy: 0.1667 - val_loss: 0.7
Epoch 15/20
105/105 496s 4s/step - accuracy: 0.3323 - loss: 0.7021 - val_accuracy: 0.8333 - val_loss: 0.6
Epoch 16/20
105/105 469s 4s/step - accuracy: 0.5952 - loss: 0.6783 - val_accuracy: 0.1667 - val_loss: 0.7
Epoch 17/20
105/105 463s 4s/step - accuracy: 0.5938 - loss: 0.6816 - val_accuracy: 0.1667 - val_loss: 0.7
Epoch 18/20
105/105 509s 4s/step - accuracy: 0.4427 - loss: 0.6943 - val_accuracy: 0.8333 - val_loss: 0.6
Epoch 19/20
105/105 494s 4s/step - accuracy: 0.5890 - loss: 0.6959 - val_accuracy: 0.8333 - val_loss: 0.6
Epoch 20/20
105/105 509s 4s/step - accuracy: 0.5189 - loss: 0.6957 - val_accuracy: 0.1667 - val_loss: 0.6
27/27 - 89s - 3s/step

EfficientNetB0 Model Accuracy: 0.8333333333333334

EfficientNetB0 Classification Report:

	precision	recall	f1-score	support
0	0.83	1.00	0.91	700
1	0.00	0.00	0.00	140
	- - -	- - -	- - -	- - -

accuracy			0.83	840
macro avg	0.42	0.50	0.45	840
weighted avg	0.69	0.83	0.76	840

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
1 # InceptionV3 Model
2 from tensorflow.keras.layers import GlobalAveragePooling2D, Dropout, Dense
3 from tensorflow.keras.models import Model
4 from sklearn.metrics import accuracy_score, classification_report
5 from tensorflow.keras.applications import InceptionV3
6
7 inception_model = InceptionV3(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
8 inception_model.trainable = False
9
10 x = inception_model.output
11 x = GlobalAveragePooling2D()(x)
12 dropout_layer = Dropout(0.5)(x)
13 output = Dense(1, activation='sigmoid')(dropout_layer)
14
15 final_inception_model = Model(inputs=inception_model.input, outputs=output)
16
17 final_inception_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
18 inception_history = final_inception_model.fit(
19     X_train_res_images, y_train_res, epochs=20, class_weight=class_weights
20 )
21
22 inception_predictions = final_inception_model.predict(
23     val_data, steps=int(np.ceil(val_data.samples / val_data.batch_size)), verbose=2
24 )
25 inception_predictions = (inception_predictions > 0.6).astype(int)
26
27 print("InceptionV3 Model Accuracy:", accuracy_score(val_data.classes, inception_predictions))
28 print("InceptionV3 Classification Report:\n", classification_report(val_data.classes, inception_predictions))
~~
```

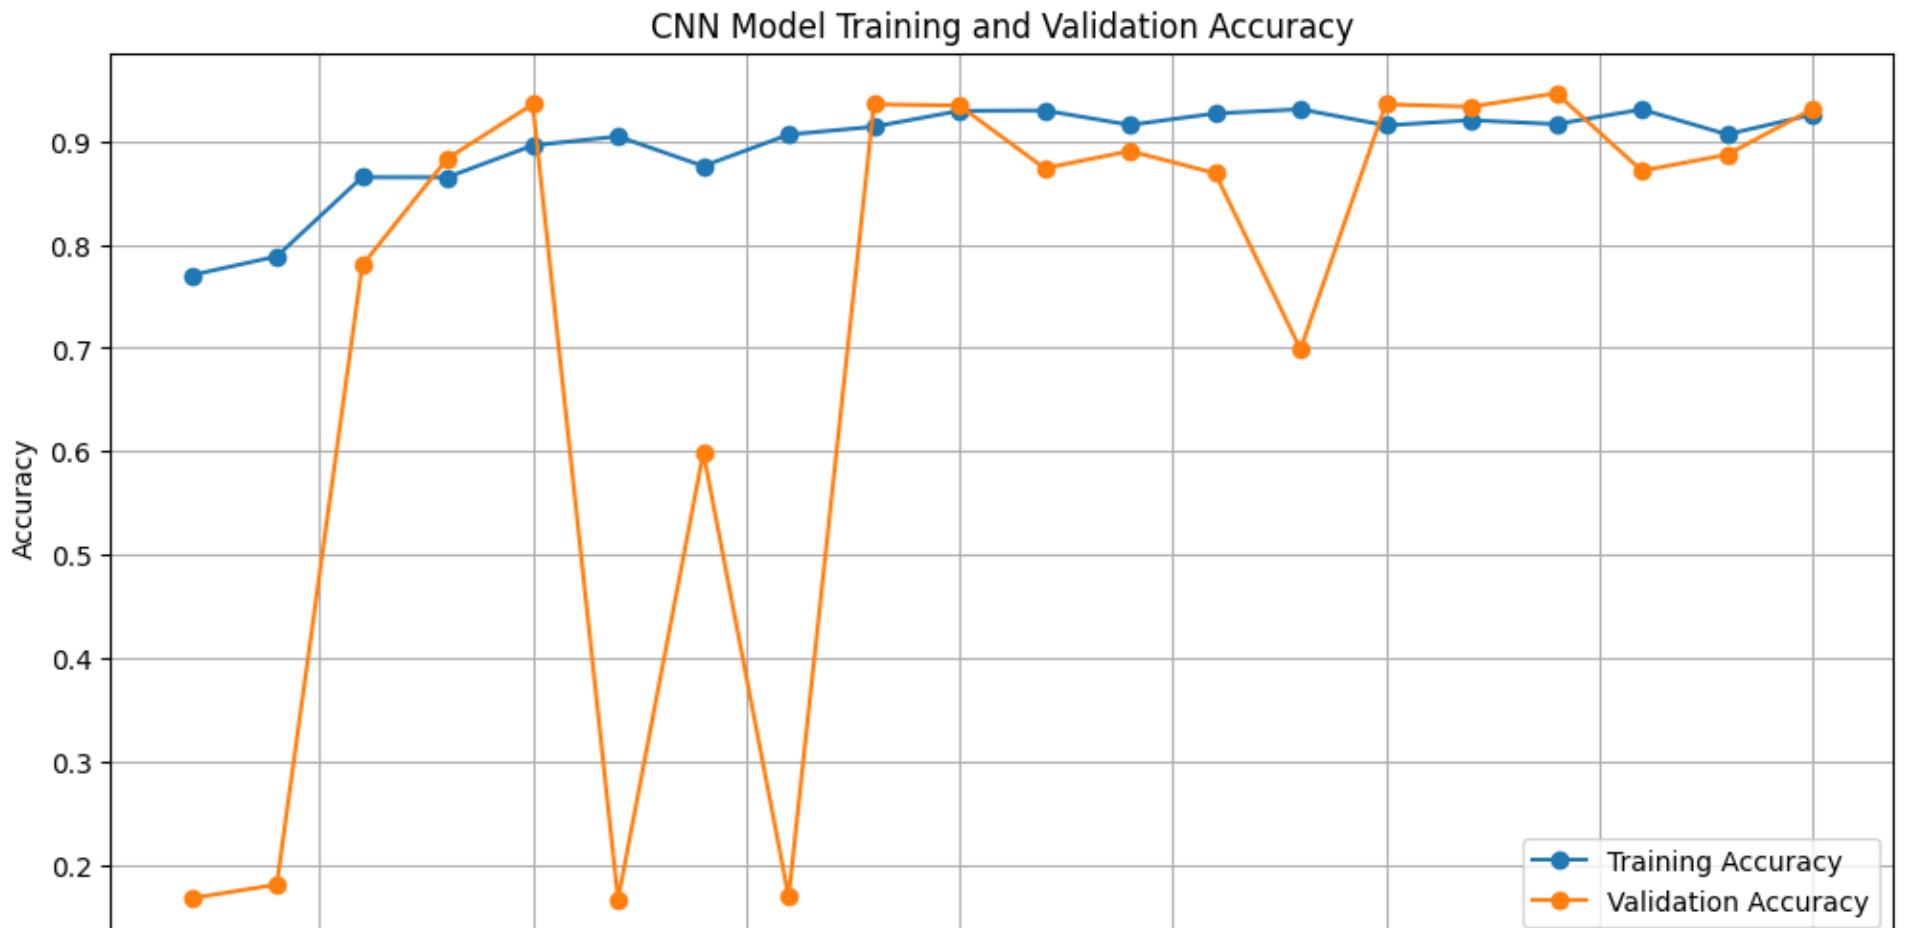
29

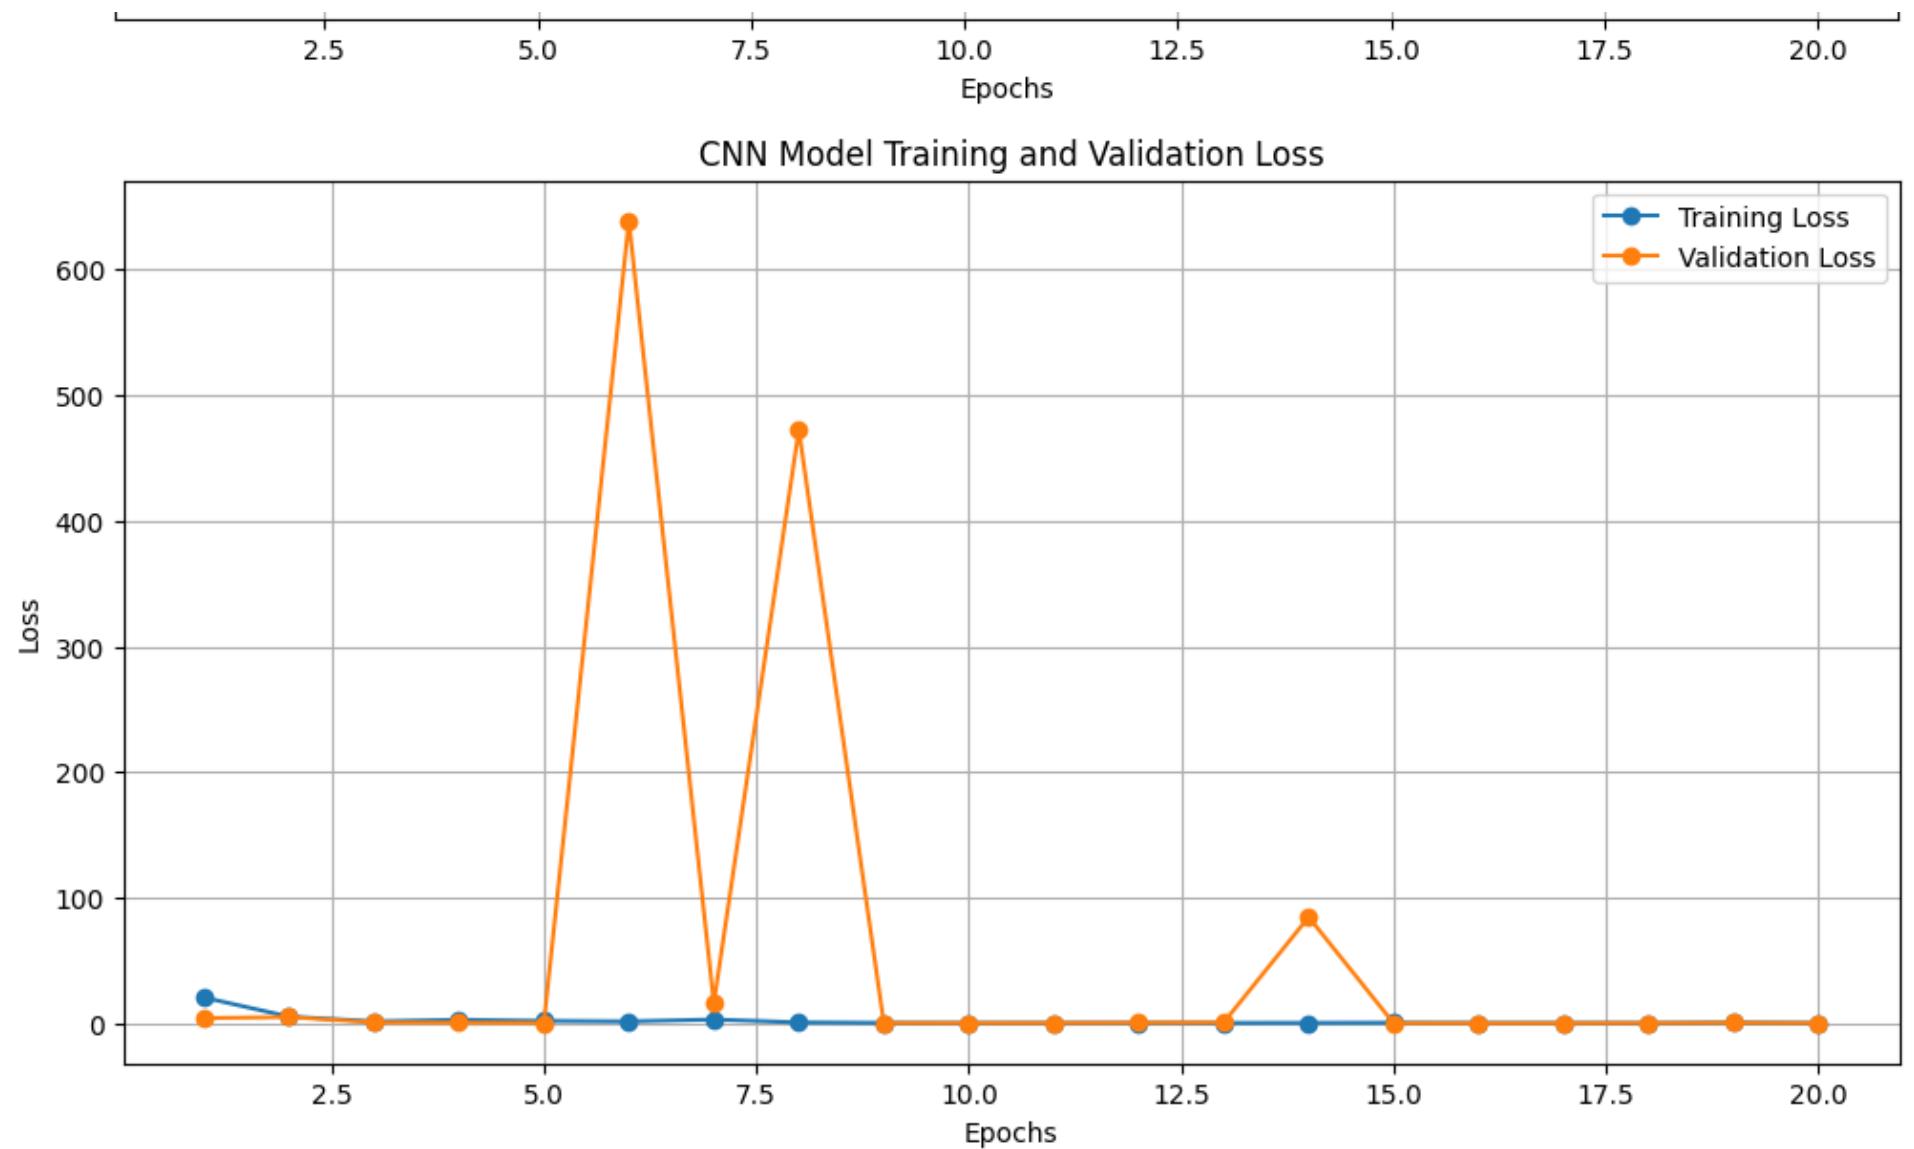
```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception\_v3/inception\_v3\_weights\_tf\_dim\_ordering\_tf\_kernels.h5
87910968/87910968 0s 0us/step
Epoch 1/20
105/105 681s 6s/step - accuracy: 0.7082 - loss: 0.5768 - val_accuracy: 0.9060 - val_loss: 0.2
Epoch 2/20
105/105 662s 6s/step - accuracy: 0.8837 - loss: 0.3111 - val_accuracy: 0.9155 - val_loss: 0.2
Epoch 3/20
105/105 705s 6s/step - accuracy: 0.9083 - loss: 0.2428 - val_accuracy: 0.9381 - val_loss: 0.2
Epoch 4/20
105/105 667s 6s/step - accuracy: 0.9091 - loss: 0.2317 - val_accuracy: 0.9250 - val_loss: 0.1
Epoch 5/20
105/105 689s 6s/step - accuracy: 0.9201 - loss: 0.2431 - val_accuracy: 0.9190 - val_loss: 0.1
Epoch 6/20
105/105 688s 6s/step - accuracy: 0.9207 - loss: 0.2189 - val_accuracy: 0.9464 - val_loss: 0.1
Epoch 7/20
105/105 745s 7s/step - accuracy: 0.9281 - loss: 0.2102 - val_accuracy: 0.9274 - val_loss: 0.2
Epoch 8/20
105/105 668s 6s/step - accuracy: 0.9189 - loss: 0.2123 - val_accuracy: 0.9488 - val_loss: 0.1
Epoch 9/20
105/105 694s 6s/step - accuracy: 0.9314 - loss: 0.1949 - val_accuracy: 0.9381 - val_loss: 0.1
Epoch 10/20
105/105 673s 6s/step - accuracy: 0.9443 - loss: 0.1692 - val_accuracy: 0.9452 - val_loss: 0.1
Epoch 11/20
105/105 673s 6s/step - accuracy: 0.9280 - loss: 0.1842 - val_accuracy: 0.9381 - val_loss: 0.1
Epoch 12/20
105/105 662s 6s/step - accuracy: 0.9227 - loss: 0.2294 - val_accuracy: 0.9405 - val_loss: 0.1
Epoch 13/20
105/105 662s 6s/step - accuracy: 0.9301 - loss: 0.1801 - val_accuracy: 0.9369 - val_loss: 0.1
Epoch 14/20
105/105 682s 6s/step - accuracy: 0.9446 - loss: 0.1587 - val_accuracy: 0.9429 - val_loss: 0.1
Epoch 15/20
105/105 665s 6s/step - accuracy: 0.9344 - loss: 0.1738 - val_accuracy: 0.9452 - val_loss: 0.1
Epoch 16/20
105/105 697s 6s/step - accuracy: 0.9355 - loss: 0.1747 - val_accuracy: 0.9440 - val_loss: 0.1
Epoch 17/20
105/105 687s 6s/step - accuracy: 0.9364 - loss: 0.1885 - val_accuracy: 0.9440 - val_loss: 0.1
Epoch 18/20
105/105 719s 6s/step - accuracy: 0.9165 - loss: 0.2286 - val_accuracy: 0.9393 - val_loss: 0.1
Epoch 19/20
```

```
Epoch 19/20
105/105 702s 6s/step - accuracy: 0.9392 - loss: 0.1547 - val_accuracy: 0.9417 - val_loss: 0.1
Epoch 20/20
105/105 681s 6s/step - accuracy: 0.9362 - loss: 0.1700 - val_accuracy: 0.9417 - val_loss: 0.1
27/27 - 129s - 5s/step
InceptionV3 Model Accuracy: 0.9464285714285714
InceptionV3 Classification Report:
precision    recall   f1-score   support
          0       0.94      0.99      0.97      700
          1       0.96      0.71      0.81      140
accuracy                           0.95      840
macro avg       0.95      0.85      0.89      840
weighted avg    0.95      0.95      0.94      840
```

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
4
5 plt.figure(figsize=(12, 6))
6 plt.plot(cnn_history.history['accuracy'], label="Training Accuracy", marker='o')
7 plt.plot(cnn_history.history['val_accuracy'], label="Validation Accuracy", marker='o')
8 plt.title("CNN Model Training and Validation Accuracy")
9 plt.xlabel("Epochs")
10 plt.ylabel("Accuracy")
11 plt.legend()
12 plt.grid(True)
13 plt.show()
14
15
16 plt.figure(figsize=(12, 6))
17 plt.plot(cnn_history.history['loss'], label="Training Loss", marker='o')
18 plt.plot(cnn_history.history['val_loss'], label="Validation Loss", marker='o')
19 plt.title("CNN Model Training and Validation Loss")
20 plt.xlabel("Epochs")
21 plt.ylabel("Loss")
```

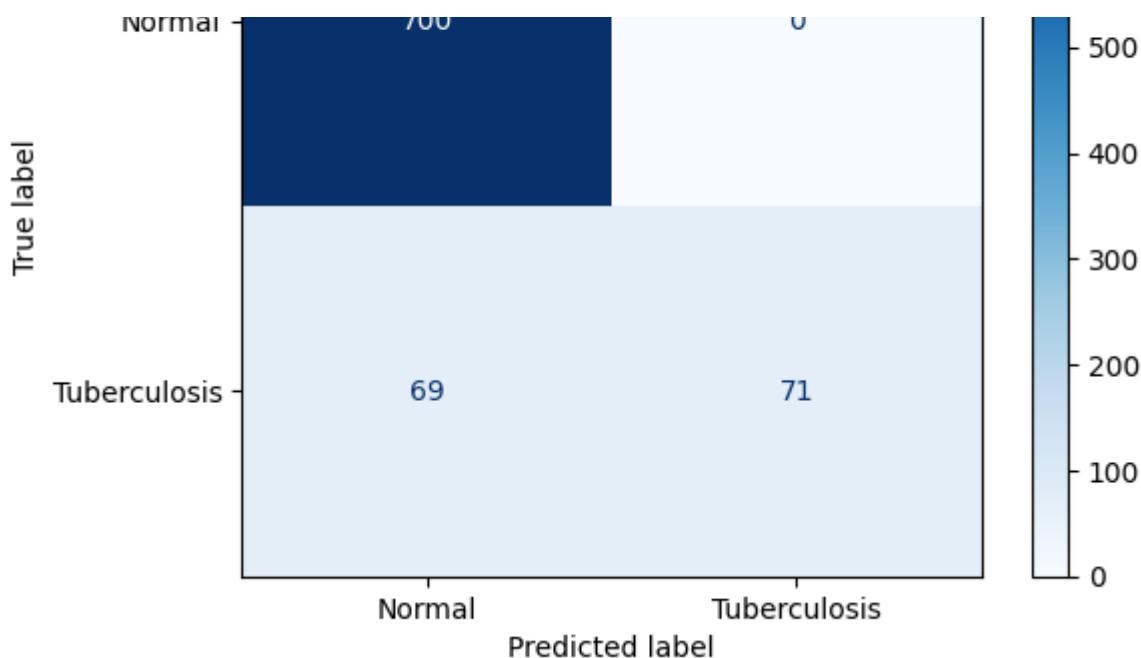
```
22 plt.legend()  
23 plt.grid(True)  
24 plt.show()  
25  
26 cm = confusion_matrix(val_data.classes, cnn_predictions)  
27 cm = confusion_matrix(y_true, y_pred)  
28 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Normal', 'Tuberculosis'])  
29 disp.plot(cmap='Blues', values_format='d')  
30 plt.title("CNN Model Confusion Matrix")  
31 plt.show()  
32
```





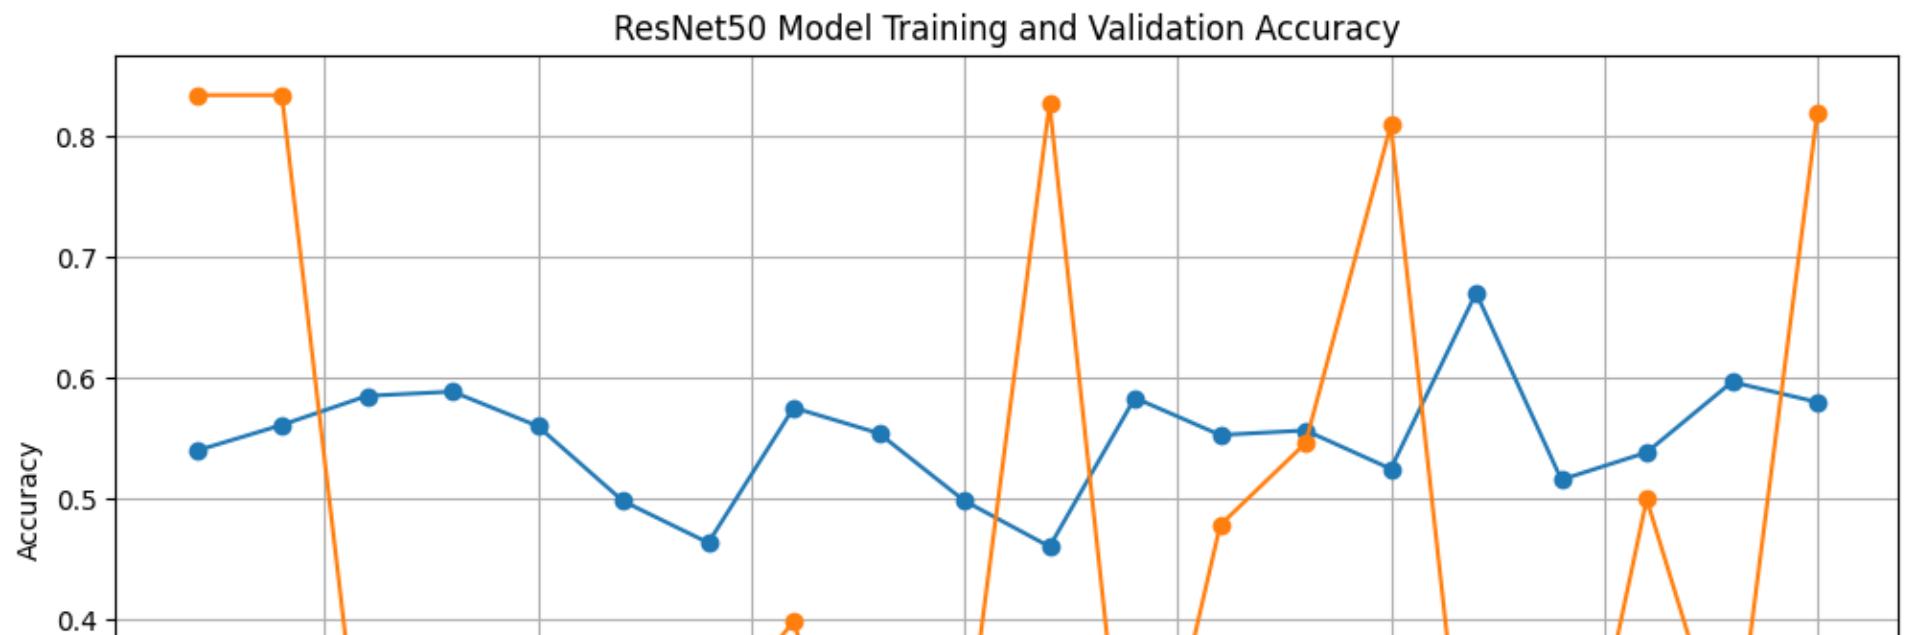
CNN Model Confusion Matrix

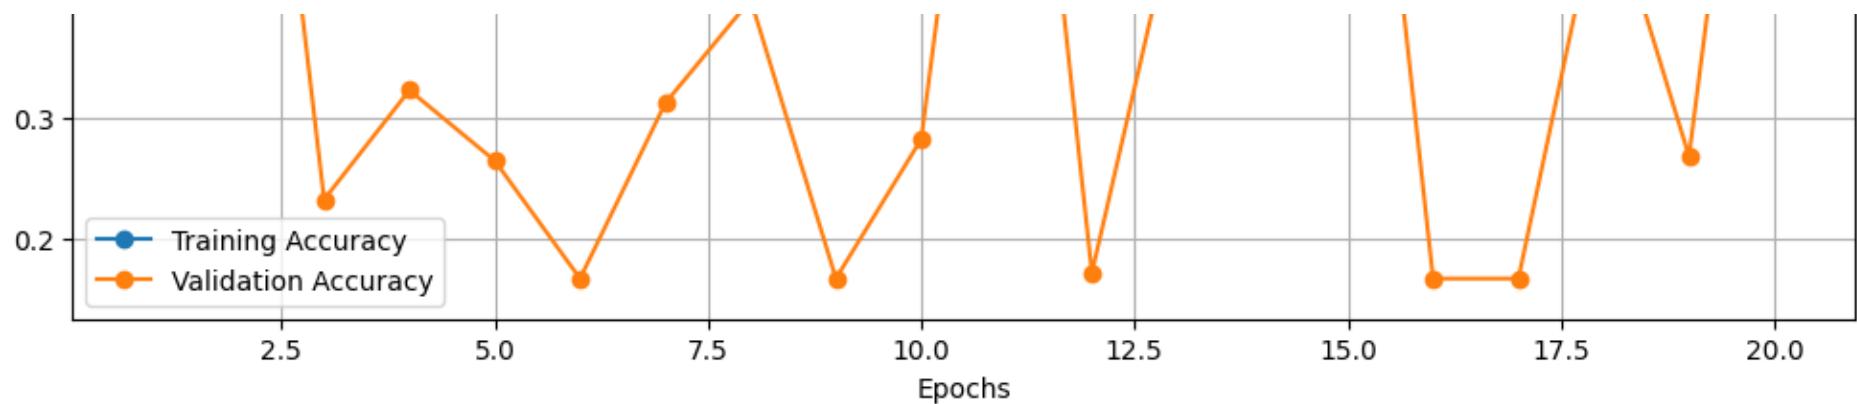




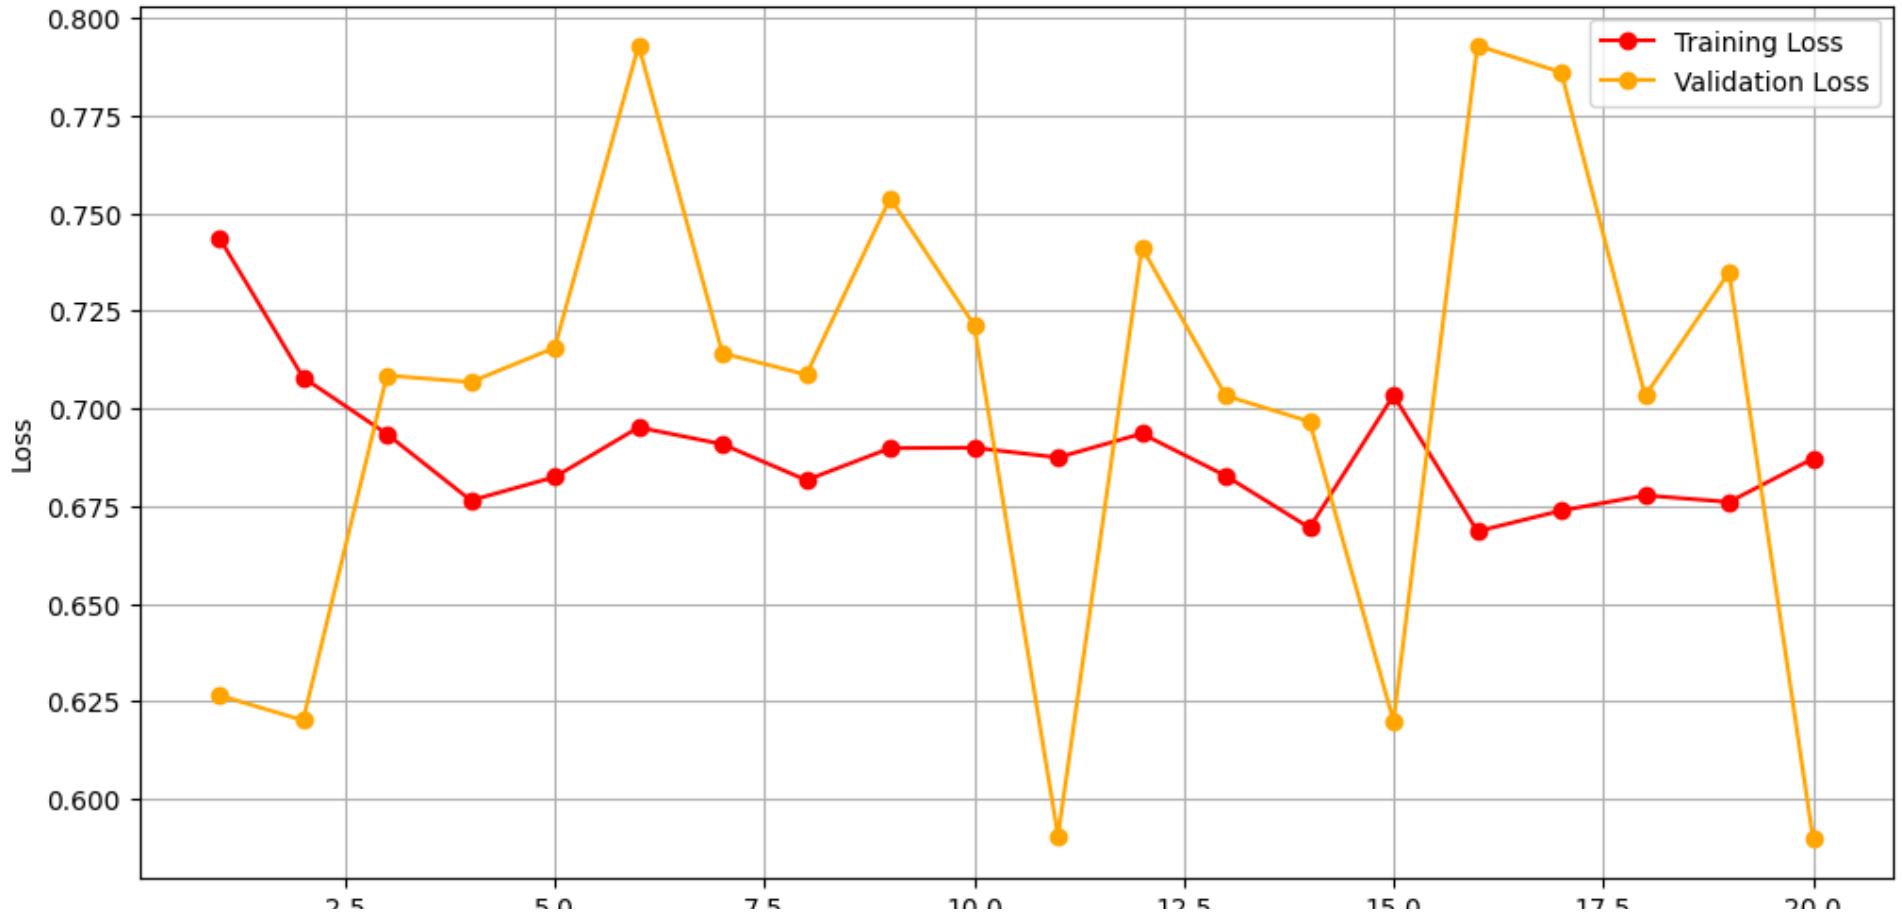
```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
4
5
6 plt.figure(figsize=(12, 6))
7 plt.plot(resnet_history.history['accuracy'], label='Training Accuracy', marker='o')
8 plt.plot(resnet_history.history['val_accuracy'], label='Validation Accuracy', marker='o')
9 plt.title('ResNet50 Model Training and Validation Accuracy')
10 plt.xlabel('Epochs')
11 plt.ylabel('Accuracy')
12 plt.legend()
13 plt.grid(True)
14 plt.show()
15
```

```
16  
17 plt.figure(figsize=(12, 6))  
18 plt.plot(resnet_history.history['loss'], label='Training Loss', marker='o', color='red')  
19 plt.plot(resnet_history.history['val_loss'], label='Validation Loss', marker='o', color='orange')  
20 plt.title('ResNet50 Model Training and Validation Loss')  
21 plt.xlabel('Epochs')  
22 plt.ylabel('Loss')  
23 plt.legend()  
24 plt.grid(True)  
25 plt.show()  
26  
27  
28 cm = confusion_matrix(val_data.classes, resnet_predictions)  
29 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Normal', 'Tuberculosis'])  
30 disp.plot(cmap='Blues')  
31 plt.title('ResNet50 Model Confusion Matrix')  
32 plt.show()  
33  
34
```



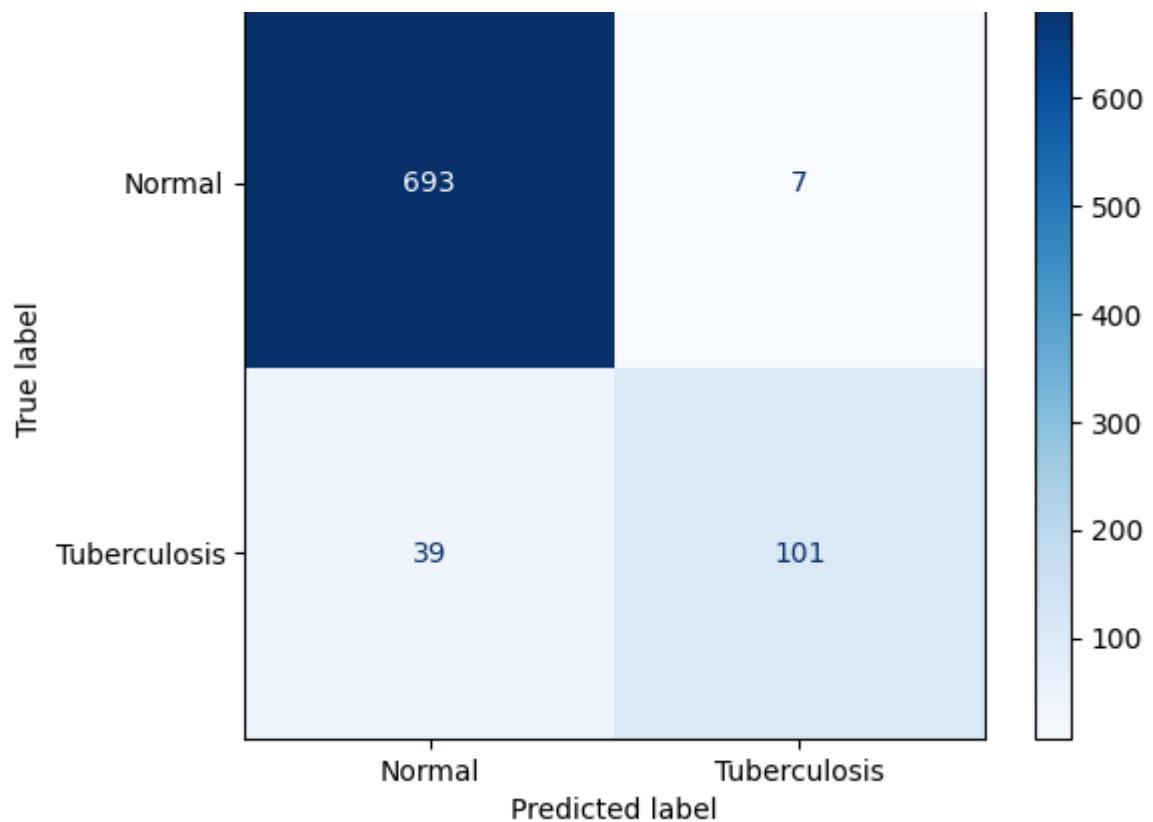


ResNet50 Model Training and Validation Loss

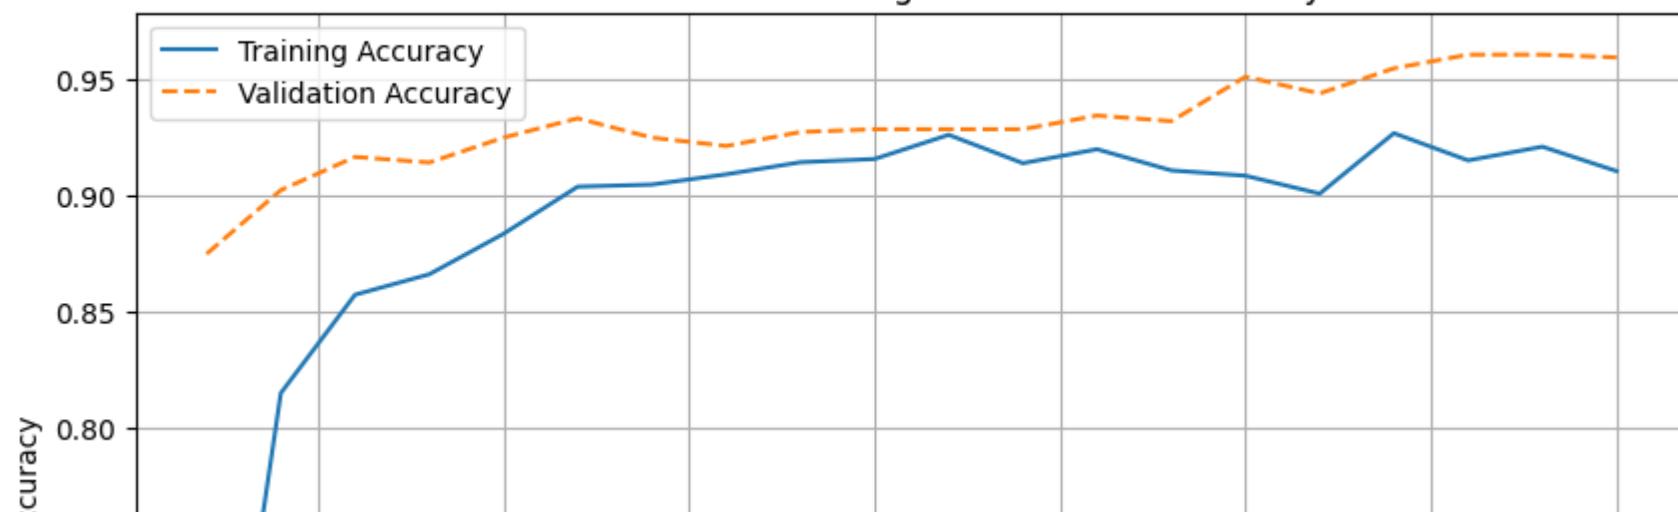


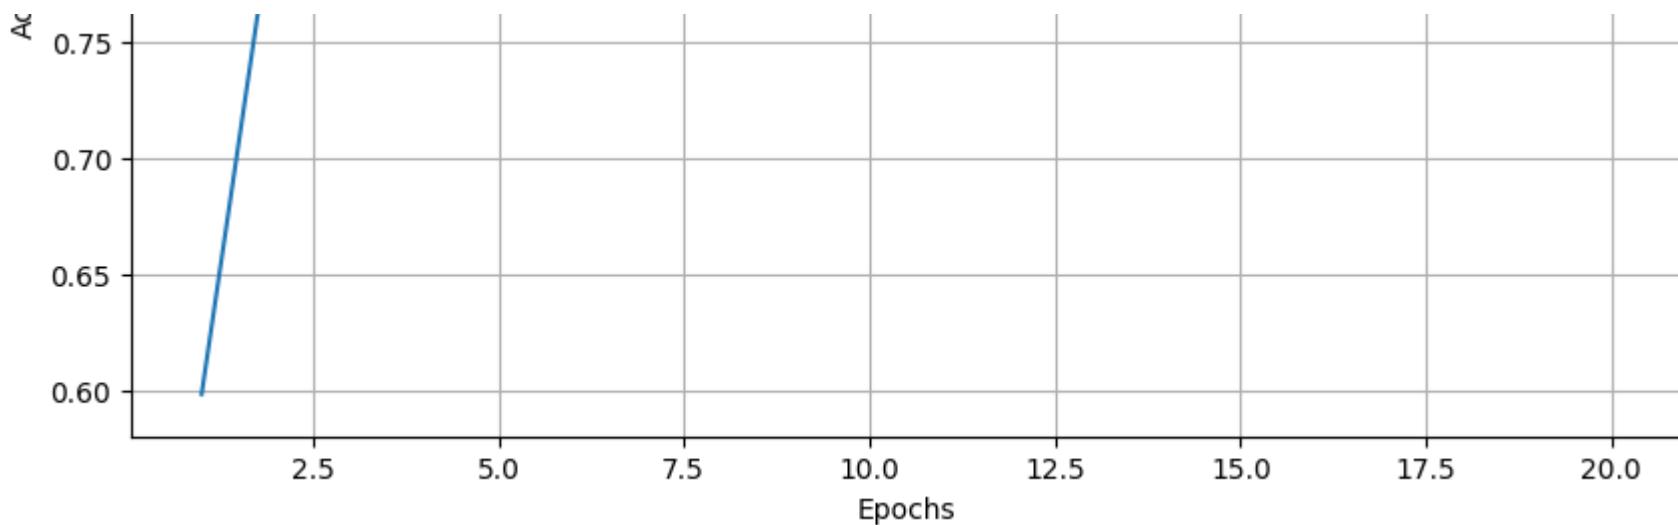
```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
4
5 cm = confusion_matrix(val_data.classes, cnn_predictions)
6 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Normal", "Tuberculosis"])
7
8
9 disp.plot(cmap=plt.cm.Blues)
10 plt.title("DenseNet Model Confusion Matrix")
11 plt.show()
12
13
14 plt.figure(figsize=(10, 6))
15 plt.plot(densenet_history.history['accuracy'], label="Training Accuracy")
16 plt.plot(densenet_history.history['val_accuracy'], label="Validation Accuracy", linestyle="--")
17 plt.xlabel("Epochs")
18 plt.ylabel("Accuracy")
19 plt.title("DenseNet Model Training and Validation Accuracy")
20 plt.legend()
21 plt.grid()
22 plt.show()
23
24
25 plt.figure(figsize=(10, 6))
26 plt.plot(densenet_history.history['loss'], label="Training Loss")
27 plt.plot(densenet_history.history['val_loss'], label="Validation Loss", linestyle="--")
28 plt.xlabel("Epochs")
29 plt.ylabel("Loss")
30 plt.title("DenseNet Model Training and Validation Loss")
31 plt.legend()
32 plt.grid()
33 plt.show()
34
```

DenseNet Model Confusion Matrix



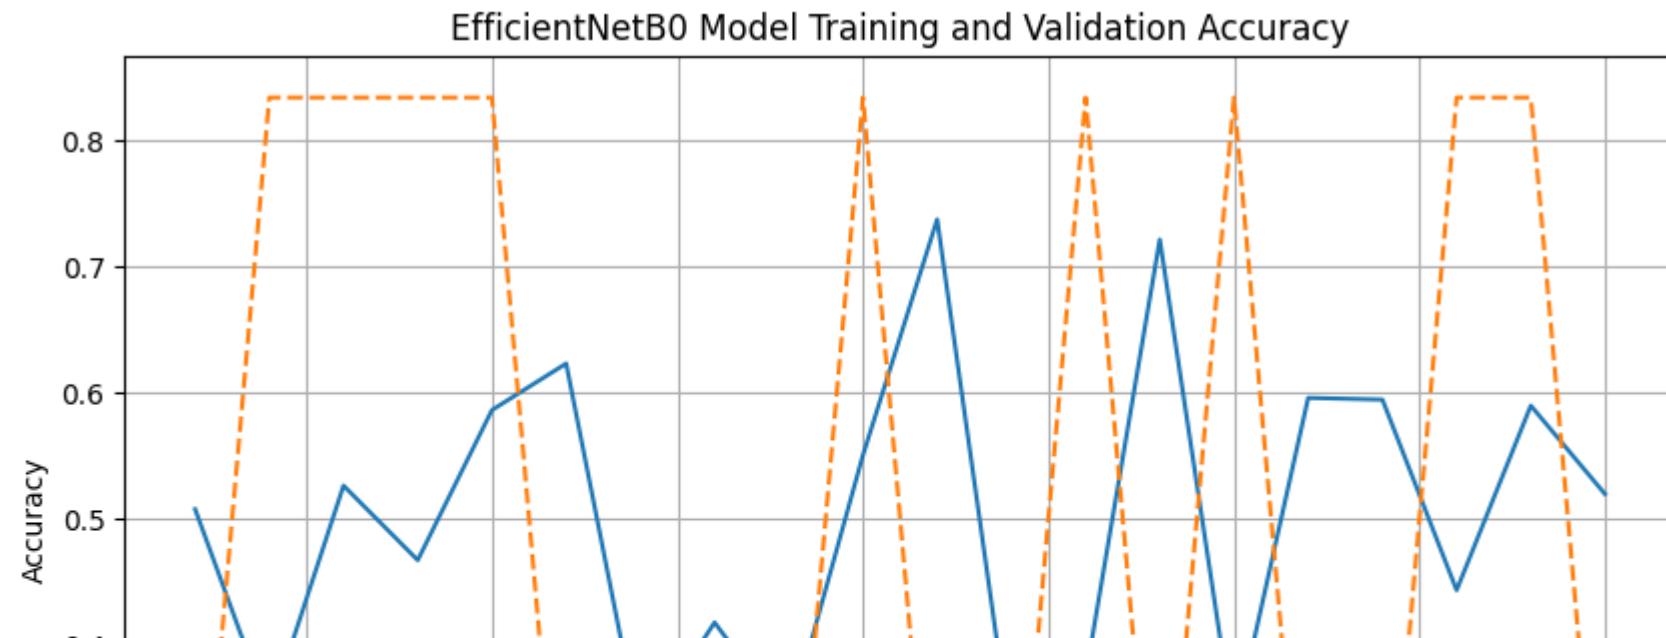
DenseNet Model Training and Validation Accuracy

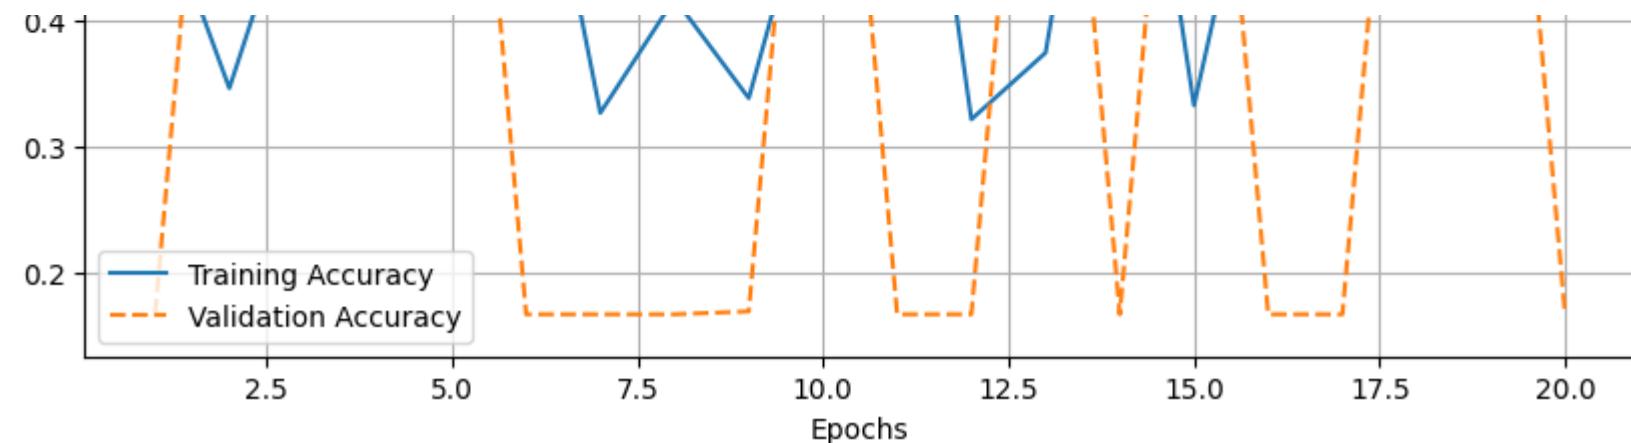




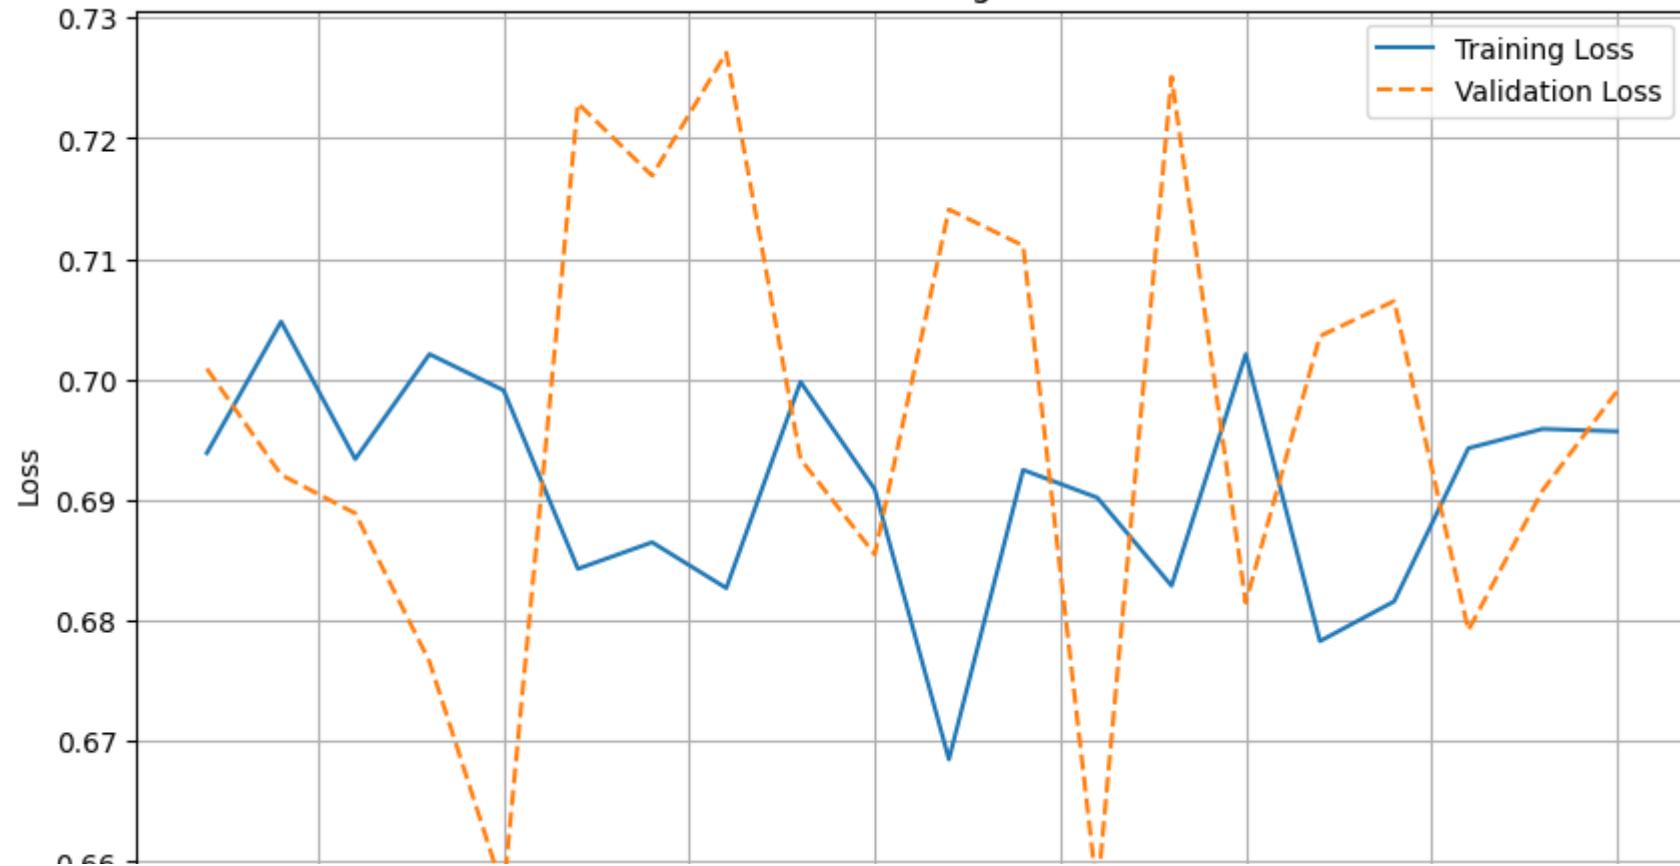
```
1 import matplotlib.pyplot as plt
2
3
4
5 plt.figure(figsize=(10, 6))
6 plt.plot(efficientnet_history.history['accuracy'], label="Training Accuracy")
7 plt.plot(efficientnet_history.history['val_accuracy'], '--', label="Validation Accuracy")
8 plt.title("EfficientNetB0 Model Training and Validation Accuracy")
9 plt.xlabel("Epochs")
10 plt.ylabel("Accuracy")
11 plt.legend()
12 plt.grid(True)
13 plt.show()
14
15
16
17
18 plt.figure(figsize=(10, 6))
19 plt.plot(efficientnet_history.history['loss'], label="Training Loss")
20 plt.plot(efficientnet_history.history['val_loss'], '--', label="Validation Loss")
21 plt.title("EfficientNetB0 Model Training and Validation Loss")
22 plt.xlabel("Epochs")
```

```
23 plt.ylabel("Loss")
24 plt.legend()
25 plt.grid(True)
26 plt.show()
27 import numpy as np
28 import seaborn as sns
29 import matplotlib.pyplot as plt
30
31
32 cm = confusion_matrix(val_data.classes, efficientnet_predictions)
33
34
35 plt.figure(figsize=(6, 5))
36 sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["Normal", "Tuberculosis"], yticklabels=["Normal", 'Tuberculosis"])
37 plt.title("Matrix for EfficientNetB0")
38 plt.xlabel("Predicted")
39 plt.ylabel("Actual")
40 plt.show()
41
```



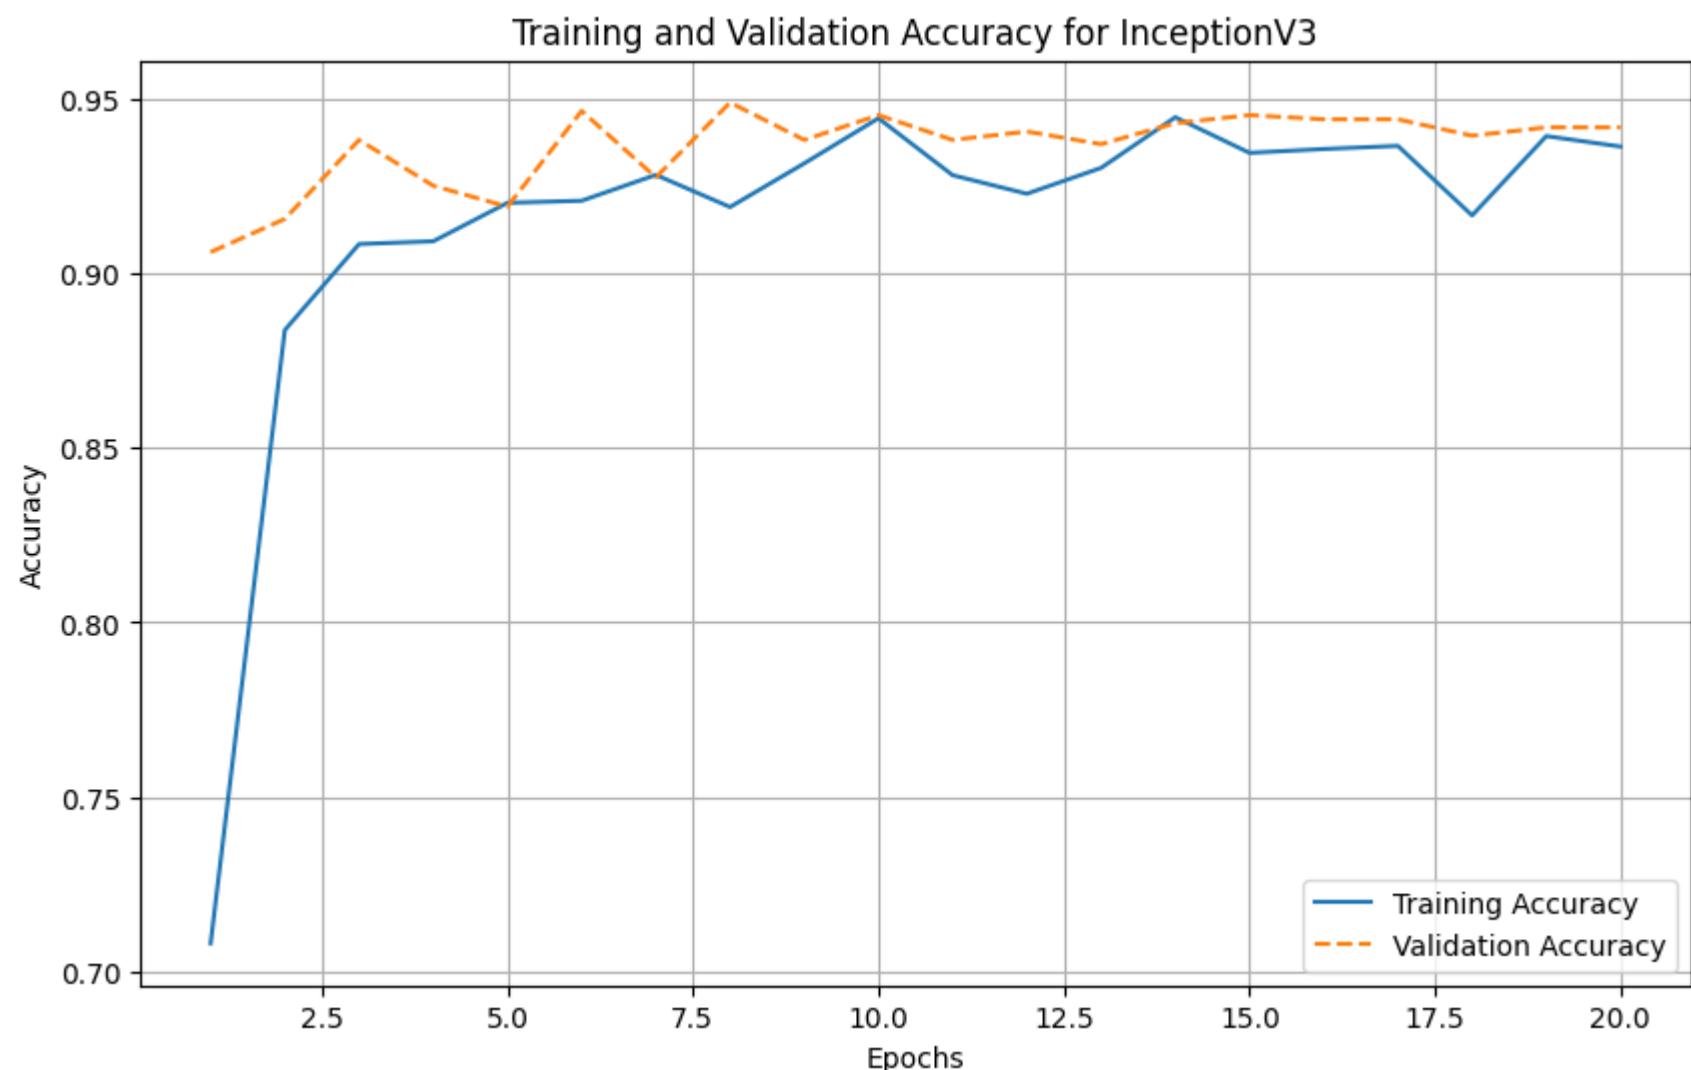


EfficientNetB0 Model Training and Validation Loss

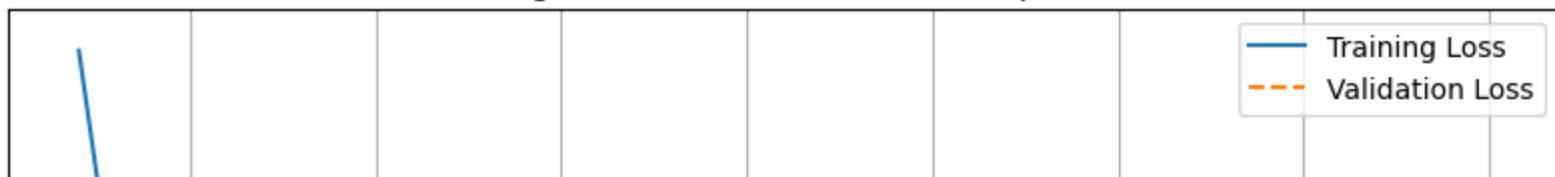


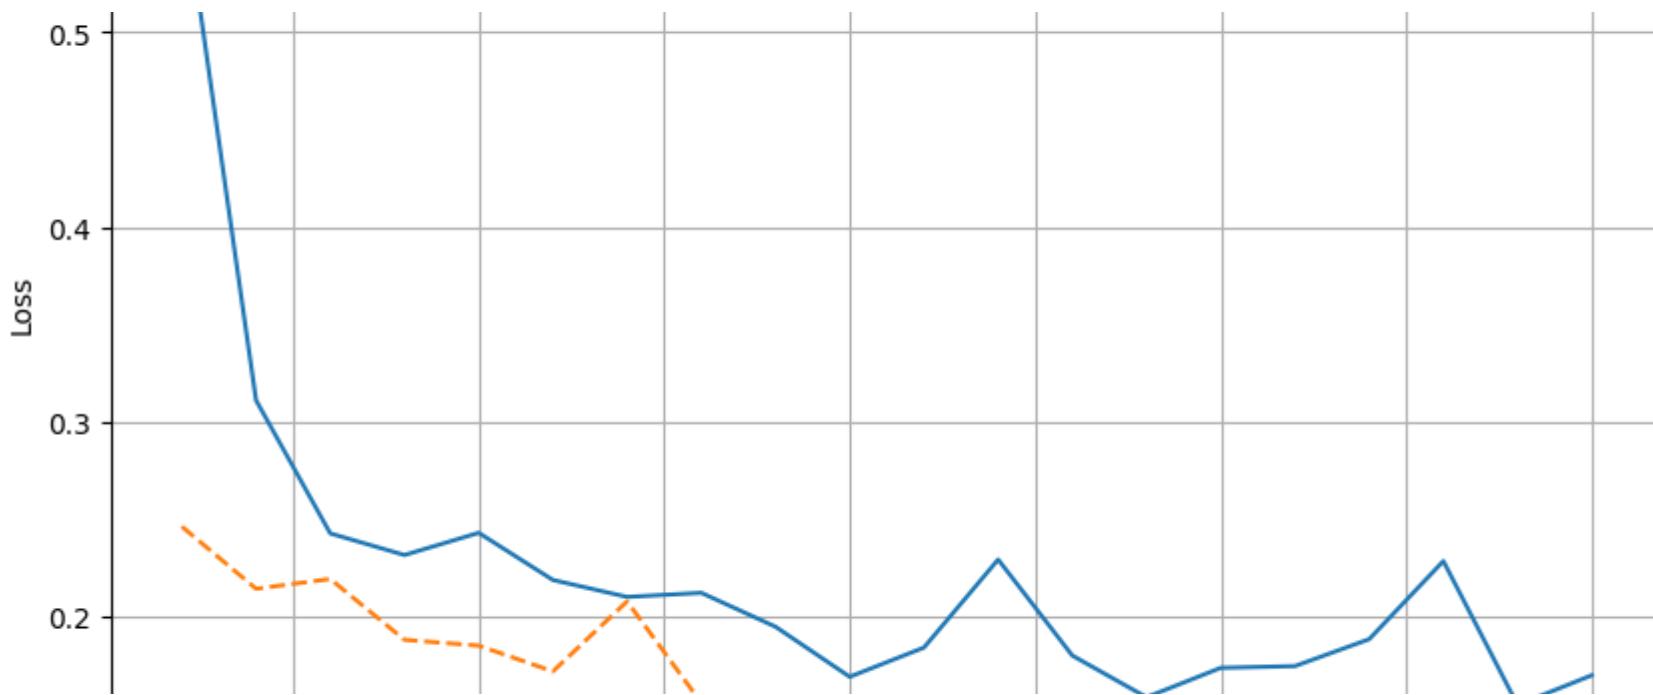
```
1 import numpy as np
2
3
4 plt.figure(figsize=(10, 6))
5 plt.plot(inception_history.history['accuracy'], label="Training Accuracy")
6 plt.plot(inception_history.history['val_accuracy'], label="Validation Accuracy", linestyle='--')
7 plt.title("Training and Validation Accuracy for InceptionV3")
8 plt.xlabel("Epochs")
9 plt.ylabel("Accuracy")
10 plt.legend()
11 plt.grid()
12 plt.show()
13
14
15 plt.figure(figsize=(10, 6))
16 plt.plot(inception_history.history['loss'], label="Training Loss")
17 plt.plot(inception_history.history['val_loss'], label="Validation Loss", linestyle='--')
18 plt.title("Training and Validation Loss for InceptionV3")
19 plt.xlabel("Epochs")
20 plt.ylabel("Loss")
21 plt.legend()
22 plt.grid()
23 plt.show()
24 from sklearn.metrics import confusion_matrix
25 import seaborn as sns
26
27 cm = confusion_matrix(val_data.classes, inception_predictions)
28
29
30 plt.figure(figsize=(6, 5))
31 sns.heatmap(cm, annot=True, fmt=".0f", cmap="Blues",
32             xticklabels=["Normal", "Tuberculosis"],
33             yticklabels=["Normal", "Tuberculosis"])
34 plt.title("Confusion Matrix for InceptionV3")
35 plt.xlabel("Predicted")
36 plt.ylabel("Actual")
37 plt.show()
```

38



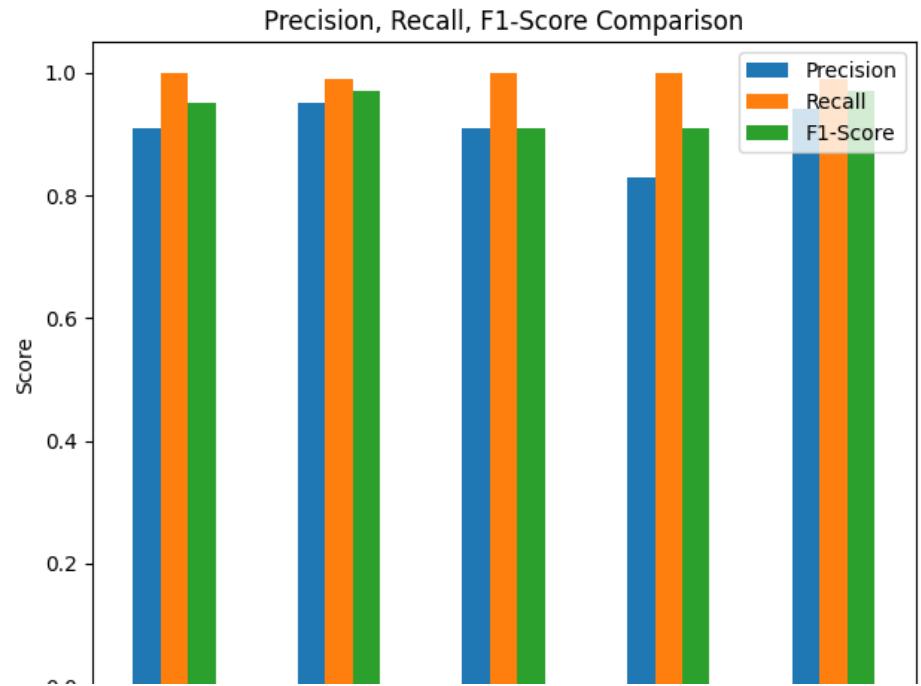
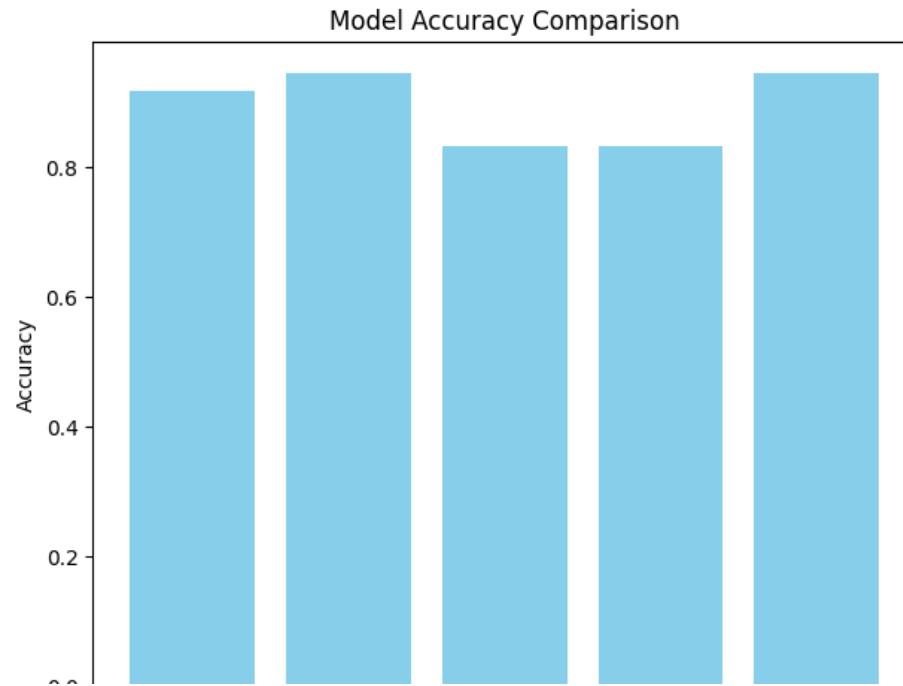
Training and Validation Loss for InceptionV3





```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3 import numpy as np
4 import pandas as pd
5
6
7 models = ['CNN', 'DenseNet', 'ResNet50', 'EfficientNetB0', 'InceptionV3']
8 accuracies = [0.917857, 0.946429, 0.833333, 0.833333, 0.946429]
9 precision = [0.91, 0.95, 0.91, 0.83, 0.94]
10 recall = [1.00, 0.99, 1.00, 1.00, 0.99]
11 f1_score = [0.95, 0.97, 0.91, 0.91, 0.97]
12
13
14 fig, axes = plt.subplots(1, 2, figsize=(14, 6))
15 fig.tight_layout(pad=5.0)
16
17
18 axes[0].bar(models, accuracies, color='skyblue')
```

```
19 axes[0].set_title('Model Accuracy Comparison')
20 axes[0].set_ylabel('Accuracy')
21
22
23 metrics_df = pd.DataFrame({
24     'Model': models,
25     'Precision': precision,
26     'Recall': recall,
27     'F1-Score': f1_score
28 })
29
30 metrics_df.set_index('Model').plot(kind='bar', ax=axes[1], color=['#1f77b4', '#ff7f0e', '#2ca02c'])
31 axes[1].set_title('Precision, Recall, F1-Score Comparison')
32 axes[1].set_ylabel('Score')
33
34
35 plt.show()
36
```



U.U CNN DenseNet ResNet50 EfficientNetB0 InceptionV3

U.U CNN DenseNet ResNet50 EfficientNetB0 InceptionV3
Model